



404 Not Found

Truong Minh Hieu
Nguyen Tri Khoa
Le Nguyen Quoc Binh

2026-10-08

ICPC Strategy

By Lukáš Poláček

This article is about my experience with ACM contests. I learned this from other people but mostly from my own and my team's mistakes.

These are the three most important facts you should bear in mind during the ACM ICPC contest:

1. It is a team competition and each team has 3 members.
2. Computer time is very expensive.
3. The beginning of the competition is the most important part of the contest.

General strategy

From these facts you can derive a very good strategy.

At the beginning of the contest, the quickest typist writes code templates and a few scripts (e.g. 'c' for compiling with the right flags). During the same time, the two others skim through the problem set in search for an easy problem (one person going from the back, one from the front). As soon as a sufficiently easy problem is located (one that can be solved within, say, 10-15 minutes), it's given to the quickest typist along with a quick description. Due to ICPC scoring it's optimal to solve easier problems first.

The other two non-coding teammates continue skimming through the problem set. If something easier than the first problem is located (solvable in 5 minutes), it takes priority. After at least one member of the team has read each problem statement, the two non-coding teammates should discuss all the problems they have just read.

What should not happen in the beginning: After reading a problem I see that I can solve it. I don't read other problem statements and solve this problem for next 2 hours without success. But there were at least 2 other easier tasks that I didn't read. That's why

reading all problems and discussion between members in the beginning is so important.

The team continues with solving problems. If someone knows how to solve a problem and nobody uses computer, he should code it. If the computer is not available, he should prepare the code on the paper.

Useful advices

You are now familiar with the general strategy. We continue with useful advices.

- If you are not sure about your solution, discuss it with your teammates. If stuck, explain problem to a teammate. Use good judgement if it's worth the interruption of him/her.
- If you have time, write code on the paper before you go the computer. It will save a lot of computer time of the whole team. You don't have to write whole code but try to focus on the most important parts of the program. For example if you are writing binary search, make sure you are using the right invariant.
- Do not debug code on the computer. Print your code and debugging output and debug on paper.
- If you are stuck on a problem, take a walk or go to the toilet. The best ideas come to mind here.
- If you keep getting WA on a problem, let it be for a while and try to solve another problem. Maybe you will get an idea how to solve it afterwards. Don't hesitate to do a complete rewrite of a solution. For most of the problems it can be done in 15 minutes.
- Is it easy to generate some large inputs, or inputs where you know the answer? If so, it may be worth doing so to test a bit more before submitting.

- When you're done with a problem, throw all papers concerning that problem onto the floor (problem set page, printouts, handwritten stuff). Saves some time searching for paper, and feels good.

- Look at the scoreboard every now and then. If there is a problem that all other teams solved, it should be easy.

- Keep track of all submissions on a sheet of paper, and keep track of who's working on what problem.

- Print early, print often. Print every time you submit.

- Don't forget the endgame strategy. When time is starting to run out you don't want all three people working on separate problems, but focus on one problem. Try to make sure that all people are still doing something useful (e.g. one person looking over the back of the coder, and another person trying to come up with tricky test cases). Knowing when to enter this mode can be difficult and in particular it takes some willpower to give up on those additional problems that one knows how to solve but just have to code up...

- So called "free submit" mode is to be used with caution. If it is not clear, its meaning is roughly "at this point, solving another problem is more important than any time penalty it can possibly cost us so let's just submit as soon as we have something that has a non-zero probability of getting accepted". Usually it is not entered until the last 30 minutes or so, but if you have a bad start with many incorrect submissions so that you already have a large penalty, it can be sometimes be entered very early (and sometimes exited if the outlook improves).

Acknowledgments

Thanks to Gunnar and Per for comments.

1	Contest
2	Mathematics
3	Numerical
4	Number theory
5	Combinatorial
6	Data structures
7	Graph
8	Geometry
9	Strings
10	Various

Contest (1)

Contest Template

```
#include <bits/stdc++.h>
using namespace std;

#define ll long long
#define ull unsigned long long
#define pii pair<int,int>
#define pll pair<long long, long long>
#define fi first
#define se second
#define all(a) (a).begin(), (a).end()
#define sz(a) (int) (a).size()
#define pb push_back
#define lwb lower_bound
#define upb upper_bound
int main() {
    ios_base::sync_with_stdio(0);
    cin.tie(0);
    cout.tie(0);
}
```

Compare Two Solutions (Stress Testing)

Description: Windows comparator: generate input, run two solutions, then compare their output files with fc. Stops at the first reported difference; leaves files for inspection. The code below is the original personal comparator with whitespace formatting only.

Usage: Compile this as a separate program and run it in the solutions' directory. Set NAME and NTEST; compile the solutions as problem.exe and problem-2.exe. Both read problem.INP; the first writes problem.OUT, the second problem.ANS (via freopen). Replace the generator under Generate tests here with valid inputs for your problem.

```
#define ll long long
using namespace std;
const string NAME = "problem";
const int NTEST = 500;
```

```
mt19937 rd(chrono::steady_clock::now().time_since_epoch().count());
#define rand rd
long long Rand(long long l, long long h) { return l + rd() % (h - l + 1); }
int main() {
    srand(time(NULL));
    for (int iT = 1; iT <= NTEST; iT++) {
        ofstream inp((NAME + ".INP").c_str());
        ll n = Rand(49, 50);
        inp << n;
        inp.close();
        system((NAME + ".exe").c_str());
        system((NAME + "-2.exe").c_str());
        if (system(("fc " + NAME + ".OUT " + NAME + ".ANS").c_str())
            != 0) {
            cout << "Test " << iT << ": WRONG!\n";
            return 0;
        }
        cout << "Test " << iT << ": CORRECT!\n";
    }
    return 0;
}
```

Troubleshooting Checklist

Pre-submit:

Write a few simple test cases if sample is not enough.
Are time limits close? If so, generate max cases.
Is the memory usage fine?
Could anything overflow?
Make sure to submit the right file.

Wrong answer:

Print your solution! Print debug output, as well.
Are you clearing all data structures between test cases?
Can your algorithm handle the whole range of input?
Read the full problem statement again.
Do you handle all corner cases correctly?
Have you understood the problem correctly?
Any uninitialized variables?
Any overflows?
Confusing N and M, i and j, etc.?
Are you sure your algorithm works?
What special cases have you not thought of?
Are you sure the STL functions you use work as you think?
Add some assertions, maybe resubmit.
Create some testcases to run your algorithm on.
Go through the algorithm for a simple case.
Go through this list again.
Explain your algorithm to a teammate.
Ask the teammate to look at your code.
Go for a small walk, e.g. to the toilet.
Is your output format correct? (including whitespace)
Rewrite your solution from the start or let a teammate do it.

Runtime error:

Have you tested all corner cases locally?
Any uninitialized variables?
Are you reading or writing outside the range of any vector?
Any assertions that might fail?
Any possible division by 0? (mod 0 for example)
Any possible infinite recursion?
Invalidated pointers or iterators?
Are you using too much memory?
Debug with resubmits (e.g. remapped signals, see Various).

Time limit exceeded:

Do you have any possible infinite loops?
What is the complexity of your algorithm?

Are you copying a lot of unnecessary data? (References)
How big is the input and output? (consider scanf)
Avoid vector, map. (use arrays/unordered_map)
What do your teammates think about your algorithm?

Memory limit exceeded:
What is the max amount of memory your algorithm should need?
Are you clearing all data structures between test cases?

Mathematics (2)

2.1 Equations

$$ax^2 + bx + c = 0 \Rightarrow x = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a}$$

The extremum is given by $x = -b/2a$.

$$\begin{aligned} ax + by &= e & x &= \frac{ed - bf}{ad - bc} \\ cx + dy &= f & y &= \frac{af - ec}{ad - bc} \end{aligned} \Rightarrow$$

In general, given an equation $Ax = b$, the solution to a variable x_i is given by

$$x_i = \frac{\det A'_i}{\det A}$$

where A'_i is A with the i 'th column replaced by b .

2.2 Recurrences

If $a_n = c_1a_{n-1} + \dots + c_ka_{n-k}$, and $r_1, \dots, r_k$ are distinct roots of $x^k - c_1x^{k-1} - \dots - c_k$, there are $d_1, \dots, d_k$ s.t.

$$a_n = d_1r_1^n + \dots + d_kr_k^n.$$

Non-distinct roots r become polynomial factors, e.g.

$$a_n = (d_1n + d_2)r^n.$$

2.3 Geometry

2.3.1 Triangles

Side lengths: a, b, c

Semiperimeter: $p = \frac{a + b + c}{2}$

Area: $A = \sqrt{p(p - a)(p - b)(p - c)}$

Circumradius: $R = \frac{abc}{4A}$

Inradius: $r = \frac{A}{p}$

Length of median (divides triangle into two equal-area triangles):
 $m_a = \frac{1}{2}\sqrt{2b^2 + 2c^2 - a^2}$

Length of bisector (divides angles in two):


$$s_a = \sqrt{bc \left[1 - \left(\frac{a}{b+c} \right)^2 \right]}$$

Law of sines: $\frac{\sin \alpha}{a} = \frac{\sin \beta}{b} = \frac{\sin \gamma}{c} = \frac{1}{2R}$

Law of cosines: $a^2 = b^2 + c^2 - 2bc \cos \alpha$

Law of tangents: $\frac{a+b}{a-b} = \frac{\tan \frac{\alpha+\beta}{2}}{\tan \frac{\alpha-\beta}{2}}$

2.3.2 Spherical coordinates



$$x = r \sin \theta \cos \phi$$
$$y = r \sin \theta \sin \phi$$
$$z = r \cos \theta$$

$$r = \sqrt{x^2 + y^2 + z^2}$$
$$\theta = \operatorname{acos}(z/\sqrt{x^2 + y^2 + z^2})$$
$$\phi = \operatorname{atan2}(y, x)$$

2.4 Derivatives/Integrals

$$\frac{d}{dx} \arcsin x = \frac{1}{\sqrt{1-x^2}}$$
$$\frac{d}{dx} \arccos x = -\frac{1}{\sqrt{1-x^2}}$$
$$\frac{d}{dx} \tan x = 1 + \tan^2 x$$
$$\frac{d}{dx} \arctan x = \frac{1}{1+x^2}$$
$$\int \tan ax = -\frac{\ln|\cos ax|}{a}$$
$$\int x \sin ax = \frac{\sin ax - ax \cos ax}{a^2}$$
$$\int e^{-x^2} = \frac{\sqrt{\pi}}{2} \operatorname{erf}(x)$$
$$\int x e^{ax} dx = \frac{e^{ax}}{a^2} (ax - 1)$$

Integration by parts:

$$\int_a^b f(x)g(x)dx = [F(x)g(x)]_a^b - \int_a^b F(x)g'(x)dx$$

2.5 Sums

$$c^a + c^{a+1} + \cdots + c^b = \frac{c^{b+1} - c^a}{c - 1}, c \neq 1$$

$$1^2 + 2^2 + 3^2 + \cdots + n^2 = \frac{n(2n+1)(n+1)}{6}$$
$$1^3 + 2^3 + 3^3 + \cdots + n^3 = \frac{n^2(n+1)^2}{4}$$
$$1^4 + 2^4 + 3^4 + \cdots + n^4 = \frac{n(n+1)(2n+1)(3n^2+3n-1)}{30}$$

2.6 Probability theory

Let X be a discrete random variable with probability $p_X(x)$ of assuming the value x . It will then have an expected value (mean) $\mu = \mathbb{E}(X) = \sum_x x p_X(x)$ and variance $\sigma^2 = V(X) = \mathbb{E}(X^2) - (\mathbb{E}(X))^2 = \sum_x (x - \mathbb{E}(X))^2 p_X(x)$ where σ is the standard deviation. If X is instead continuous it will have a probability density function $f_X(x)$ and the sums above will instead be integrals with $p_X(x)$ replaced by $f_X(x)$.

Expectation is linear:

$$\mathbb{E}(aX + bY) = a\mathbb{E}(X) + b\mathbb{E}(Y)$$

For independent X and Y ,

$$V(aX + bY) = a^2V(X) + b^2V(Y).$$

2.6.1 Discrete distributions

Binomial distribution

The number of successes in n independent yes/no experiments, each which yields success with probability p is $\operatorname{Bin}(n, p)$, $n = 1, 2, \dots$, $0 \leq p \leq 1$.

$$p(k) = \binom{n}{k} p^k (1-p)^{n-k}$$

$$\mu = np, \sigma^2 = np(1-p)$$

$\operatorname{Bin}(n, p)$ is approximately $\operatorname{Po}(np)$ for small p .

First success distribution

The number of trials needed to get the first success in independent yes/no experiments, each which yields success with probability p is $\operatorname{Fs}(p)$, $0 \leq p \leq 1$.

$$p(k) = p(1-p)^{k-1}, k = 1, 2, \dots$$

$$\mu = \frac{1}{p}, \sigma^2 = \frac{1-p}{p^2}$$

Poisson distribution

The number of events occurring in a fixed period of time t if these events occur with a known average rate κ and independently of the time since the last event is $\operatorname{Po}(\lambda)$, $\lambda = t\kappa$.

$$p(k) = e^{-\lambda} \frac{\lambda^k}{k!}, k = 0, 1, 2, \dots$$

$$\mu = \lambda, \sigma^2 = \lambda$$

2.6.2 Continuous distributions

Uniform distribution

If the probability density function is constant between a and b and 0 elsewhere it is $\operatorname{U}(a, b)$, $a < b$.

$$f(x) = \begin{cases} \frac{1}{b-a} & a < x < b \\ 0 & \text{otherwise} \end{cases}$$

$$\mu = \frac{a+b}{2}, \sigma^2 = \frac{(b-a)^2}{12}$$

Exponential distribution

The time between events in a Poisson process is $\operatorname{Exp}(\lambda)$, $\lambda > 0$.

$$f(x) = \begin{cases} \lambda e^{-\lambda x} & x \geq 0 \\ 0 & x < 0 \end{cases}$$

$$\mu = \frac{1}{\lambda}, \sigma^2 = \frac{1}{\lambda^2}$$

Normal distribution

Most real random values with mean μ and variance σ^2 are well described by $\mathcal{N}(\mu, \sigma^2)$, $\sigma > 0$.

$$f(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

If $X_1 \sim \mathcal{N}(\mu_1, \sigma_1^2)$ and $X_2 \sim \mathcal{N}(\mu_2, \sigma_2^2)$ then

$$aX_1 + bX_2 + c \sim \mathcal{N}(\mu_1 + \mu_2 + c, a^2\sigma_1^2 + b^2\sigma_2^2)$$

Numerical (3)

3.1 Polynomials and recurrences

Polynomial

Description: Polynomial coefficients `coeff[i]` multiply x^i (constant term first). Evaluation uses Horner's rule. `diff` changes the polynomial to its derivative. `divroot(r)` divides by $x-r$ and discards the remainder, so normally call it only for a known root. Use a nonempty coefficient vector and trim zero leading coefficients before root finding; do not differentiate an empty polynomial.

Usage: `Poly p{{2,-3,1}}; double value=p(3); // 2`
`p.diff(); // coefficients {-3,2}`

```
struct Poly {
    vector<double> coeff;
    double operator()(double x) const {
        double value = 0;
        for (int i = sz(coeff); i--;) (value *= x) += coeff[i];
        return value;
    }
    void diff() {
        for (int i = 1; i < (sz(coeff)); ++i) coeff[i - 1] = i *
            coeff[i];
        coeff.pop_back();
    }
    void divroot(double root) {
        double carry = coeff.back(), oldCoeff;
        coeff.back() = 0;
        for (int i = sz(coeff) - 1; i--;)
            oldCoeff = coeff[i], coeff[i] = coeff[i + 1] * root +
                carry, carry = oldCoeff;
        coeff.pop_back();
    }
};
```

Polynomial Roots

Description: Finds the real roots to a polynomial. Uses derivative roots to split the line into monotone pieces, then bisects sign changes. Use a nonconstant polynomial with nonzero leading coefficient and coefficients in increasing degree order. Choose bounds covering the roots of interest; this implementation may also return candidates outside those bounds. Repeated roots without a sign change are not guaranteed and need separate handling.

Usage: auto roots=polyRoots(Poly{{2,-3,1}},-10,10);

// approximately {1,2}

Time: $\mathcal{O}(n^2 \log(1/\epsilon))$

"Polynomial.h"

```
vector<double> polyRoots(Poly poly, double minX, double maxX) {
    if (sz(poly.coeff) == 2) return {-poly.coeff[0] / poly.coeff[1]};
    vector<double> roots;
    Poly derivative = poly;
    derivative.diff();
    auto criticalPoints = polyRoots(derivative, minX, maxX);
    criticalPoints.push_back(minX - 1);
    criticalPoints.push_back(maxX + 1);
    sort(all(criticalPoints));
    for (int i = 0; i < (sz(criticalPoints) - 1); ++i) {
        double lo = criticalPoints[i], hi = criticalPoints[i + 1];
        bool positiveLeft = poly(lo) > 0;
        if (positiveLeft ^ (poly(hi) > 0)) {
            for (int it = 0; it < (60); ++it) { // while (h - l > 1e-8)
                double mid = (lo + hi) / 2, value = poly(mid);
                if ((value <= 0) ^ positiveLeft)
                    lo = mid;
                else
                    hi = mid;
            }
            roots.push_back((lo + hi) / 2);
        }
    }
    return roots;
}
```

Polynomial Interpolation

Description: Given n points $(x[i], y[i])$, computes an $n-1$ -degree polynomial p that passes through them: $p(x) = a[0] * x^0 + ... + a[n-1] * x^{n-1}$. For numerical precision, pick $x[k] = c * \cos(k/(n-1) * \pi), k = 0...n-1$. Use distinct x coordinates and n matching x/y values. Returns n coefficients in increasing powers; result[0] is the constant term. Floating interpolation can be unstable for large n or poorly spaced x . Use small well-scaled data, or a separate modular method for exact modular answers.

Usage: auto coeff=interpolate({0,1,2},{1,4,9},3);

// approximately {1,2,1}, representing $(x+1)^2$

Time: $\mathcal{O}(n^2)$

```
typedef vector<double> vd;
vd interpolate(vd x, vd y, int n) {
    vd coeff(n), basis(n);
    for (int k = 0; k < (n - 1); ++k)
        for (int i = k + 1; i < (n); ++i) y[i] = (y[i] - y[k]) / (x[i] - x[k]);
    double previous = 0;
    basis[0] = 1;
    for (int k = 0; k < (n); ++k)
        for (int i = 0; i < (n); ++i) {
            coeff[i] += y[k] * basis[i];
            swap(previous, basis[i]);
            basis[i] -= previous * x[k];
        }
    return coeff;
}
```

Linear Recurrence

Description: Generates the k 'th term of an n -order linear recurrence $S[i] = \sum_j S[i-j-1]tr[j]$, given $S[0... \geq n-1]$ and $tr[0...n-1]$. Faster than matrix multiplication. Useful together with Berlekamp-Massey. Use when recurrence order is small but the requested index is huge. initial contains exactly n initial terms; recurrence[j] multiplies $S[i-j-1]$. index is 0-based and nonnegative. Supply the global mod used by the code; the hidden value 5 is only for repository tests. Products must fit ll.

Usage: ll fib=linearRec({0,1},{1,1},10);

// 55 modulo the configured mod

Time: $\mathcal{O}(n^2 \log k)$

```
typedef vector<ll> Poly;
ll linearRec(Poly initial, Poly recurrence, ll index) {
    int n = sz(recurrence);

    auto combine = [&](Poly a, Poly b) {
        Poly result(n * 2 + 1);
        for (int i = 0; i < (n + 1); ++i)
            for (int j = 0; j < (n + 1); ++j) result[i + j] = (result[i + j] + a[i] * b[j]) % mod;
        for (int i = 2 * n; i > n; --i)
            for (int j = 0; j < (n); ++j)
                result[i - 1 - j] = (result[i - 1 - j] + result[i] * recurrence[j]) % mod;
        result.resize(n + 1);
        return result;
    };

    Poly weights(n + 1), power(weights);
    weights[0] = power[1] = 1;

    for (++index; index; index /= 2) {
        if (index % 2) weights = combine(weights, power);
        power = combine(power, power);
    }

    ll result = 0;
    for (int i = 0; i < (n); ++i) result = (result + weights[i + 1] * initial[i]) % mod;
    return result;
}
```

3.2 Matrices

Determinant

Description: Calculates determinant of a matrix. Destroys the matrix. Input must be square. Row swaps change the determinant sign; elimination modifies the matrix, so copy it first if needed. The fixed pivot tolerance may classify very small determinants as zero; scale input when numerical precision matters.

Usage: vector<vector<double>> a={{1,2},{3,4}};

double determinant=det(a); // -2; a is modified

Time: $\mathcal{O}(N^3)$

```
double det(vector<vector<double>> &matrix) {
    int n = sz(matrix);
    double determinant = 1;
    for (int i = 0; i < (n); ++i) {
        int pivotRow = i;
        for (int j = i + 1; j < (n); ++j)
            if (fabs(matrix[j][i]) > fabs(matrix[pivotRow][i]))
                pivotRow = j;
        if (i != pivotRow) swap(matrix[i], matrix[pivotRow]),
            determinant *= -1;
        determinant *= matrix[i][i];
        if (determinant == 0) return 0;
        for (int j = i + 1; j < (n); ++j) {
            double factor = matrix[j][i] / matrix[i][i];
            if (factor != 0)
```

```
                for (int k = i + 1; k < (n); ++k) matrix[j][k] -=
                    factor * matrix[i][k];
            }
        }
        return determinant;
    }
```

Integer Determinant

Description: Calculates determinant using modular arithmetics. Modulos can also be removed to get a pure-integer version. Set global mod and normalize entries modulo it. Euclidean row steps work with composite moduli too, unlike elimination requiring inverses. Input must be square and is destroyed. Products before reduction must fit ll; the returned determinant lies in $[0, \text{mod})$.

Usage: vector<vector<ll>> a={{1,2},{3,4}};

ll determinant=det(a); // mod-2; a is modified

Time: $\mathcal{O}(N^3)$

```
const ll mod = 12345;
ll det(vector<vector<ll>> &matrix) {
    int n = sz(matrix);
    ll determinant = 1;
    for (int i = 0; i < (n); ++i) {
        for (int j = i + 1; j < (n); ++j) {
            while (matrix[j][i] != 0) { // gcd step
                ll quotient = matrix[i][i] / matrix[j][i];
                if (quotient)
                    for (int k = i; k < (n); ++k)
                        matrix[i][k] = (matrix[i][k] - matrix[j][k] *
                            quotient) % mod;
                swap(matrix[i], matrix[j]);
                determinant *= -1;
            }
        }
        determinant = determinant * matrix[i][i] % mod;
        if (!determinant) return 0;
    }
    return (determinant + mod) % mod;
}
```

Gaussian Elimination

Description: Solves $A * x = b$. If there are multiple solutions, an arbitrary one is returned. Returns rank, or -1 if no solutions. Data in A and b is lost. For n equations in m unknowns, resize solution to m before calling: its size tells the solver the number of columns. Rank m means unique solution; $0 \leq \text{rank} < m$ means free variables, chosen as zero. Read solution only if rank!=1. A and b are modified; copy them if needed.

Usage: vector<vd> a={{1,1},{1,-1}}; vd b={3,1}, x(2);

int rank=solveLinear(a,b,x); // rank=2, x={2,1}

Time: $\mathcal{O}(n^2 m)$

```
typedef vector<double> vd;
const double eps = 1e-12;
int solveLinear(vector<vd> &matrix, vd &rhs, vd &solution) {
    int n = sz(matrix), m = sz(solution), rank = 0, pivotRow,
        pivotCol;
    if (n) assert(sz(matrix[0]) == m);
    vector<int> colOrder(m);
    iota(all(colOrder), 0);

    for (int i = 0; i < (n); ++i) {
        double magnitude, pivotValue = 0;
        for (int r = i; r < (n); ++r)
            for (int c = i; c < (m); ++c)
                if ((magnitude = fabs(matrix[r][c])) > pivotValue)
                    pivotRow = r, pivotCol = c, pivotValue = magnitude;
        if (pivotValue <= eps) {
            for (int j = i; j < (n); ++j)
                if (fabs(rhs[j]) > eps) return -1;
            break;
        }
```

```
    }
    swap(matrix[i], matrix[pivotRow]);
    swap(rhs[i], rhs[pivotRow]);
    swap(colOrder[i], colOrder[pivotCol]);
    for (int j = 0; j < (n); ++j) swap(matrix[j][i], matrix[j][pivotCol]);
    pivotValue = 1 / matrix[i][i];
    for (int j = i + 1; j < (n); ++j) {
        double factor = matrix[j][i] * pivotValue;
        rhs[j] -= factor * rhs[i];
        for (int k = i + 1; k < (m); ++k) matrix[j][k] -= factor * matrix[i][k];
    }
    rank++;
}

solution.assign(m, 0);
for (int i = rank; i--;) {
    rhs[i] /= matrix[i][i];
    solution[colOrder[i]] = rhs[i];
    // for (int j = 0; j < (i); ++j) rhs[j] -= matrix[j][i] * rhs[i];
}
return rank; // (multiple solutions if rank < m)
}
```

Gaussian Elimination: Unique Values

Description: To get all uniquely determined values of x back from SolveLinear, make the following changes: This is an edit recipe for SolveLinear, not a separate callable solver. Eliminate each pivot column from all other rows, then mark variables depending on free columns as undefined. Choose an undefined marker such as NaN. A rank below m does not mean every variable is undetermined.

Usage: // Apply the changes inside solveLinear first.
// Then size x to m and call solveLinear(a,b,x) as usual.
"SolveLinear.h"

```
for (int j = 0; j < (n); ++j)
    if (j != i) // instead of for (int j = i+1; j < (n); ++j)
        // ... then at the end:
        solution.assign(m, undefined);
for (int i = 0; i < (rank); ++i) {
    for (int j = rank; j < (m); ++j)
        if (fabs(matrix[i][j]) > eps) goto fail;
    solution[colOrder[i]] = rhs[i] / matrix[i][i];
fail:;
}
```

Gaussian Elimination over GF(2)

Description: Solves $Ax = b$ over $\mathbb{F}_2$. If there are multiple solutions, one is returned arbitrarily. Returns rank, or -1 if no solutions. Destroys A and b . Use for XOR equations: each row bitset holds the first m coefficients and rhs entries are 0/1. All coefficient bits from m upward must be zero and $m \leq 1000$. Returns rank or -1; free variables are zero. Both matrix and rhs are modified. Each solution bit is that variable's assigned value.

Usage: vector<bs> a={bs(3),bs(1)}; vector<int> b={1,0}; bs x;
int rank=solveLinear(a,b,x,2); // x[0]=0, x[1]=1
Time: $\mathcal{O}(n^2m)$

```
typedef bitset<1000> bs;
int solveLinear(vector<bs> &matrix, vector<int> &rhs, bs &
    solution, int m) {
    int n = sz(matrix), rank = 0, pivotRow;
    assert(m <= sz(solution));
    vector<int> colOrder(m);
    iota(all(colOrder), 0);
    for (int i = 0; i < (n); ++i) {
        for (pivotRow = i; pivotRow < n; ++pivotRow)
            if (matrix[pivotRow].any()) break;
        if (pivotRow == n) {
            for (int j = i; j < (n); ++j)
```

```
            if (rhs[j]) return -1;
            break;
        }
        int pivotCol = (int)matrix[pivotRow]._Find_next(i - 1);
        swap(matrix[i], matrix[pivotRow]);
        swap(rhs[i], rhs[pivotRow]);
        swap(colOrder[i], colOrder[pivotCol]);
        for (int j = 0; j < (n); ++j)
            if (matrix[j][i] != matrix[j][pivotCol]) {
                matrix[j].flip(i);
                matrix[j].flip(pivotCol);
            }
        for (int j = i + 1; j < (n); ++j)
            if (matrix[j][i]) {
                rhs[j] ^= rhs[i];
                matrix[j] ^= matrix[i];
            }
        rank++;
    }

    solution = bs();
    for (int i = rank; i--;) {
        if (!rhs[i]) continue;
        solution[colOrder[i]] = 1;
        for (int j = 0; j < (i); ++j) rhs[j] ^= matrix[j][i];
    }
    return rank; // (multiple solutions if rank < m)
}
```

Matrix Inverse

Description: Invert matrix A . Returns rank; result is stored in A unless singular (rank < n). Can easily be extended to prime moduli; for prime powers, repeatedly set $A^{-1} = A^{-1}(2I - AA^{-1}) \pmod{p^k}$ where A^{-1} starts as the inverse of $A \bmod p$, and k is doubled in each step. Use a square floating matrix. Only read the overwritten matrix as the inverse when the returned rank equals its dimension; otherwise it is singular and partially modified. Row and column pivoting use a fixed tolerance. Copy the original if it is needed afterward.

Usage: vector<vector<double>>> a={{2,0},{0,4}};
int rank=matInv(a); // rank=2, a=diag(0.5,0.25)
Time: $\mathcal{O}(n^3)$

```
int matInv(vector<vector<double>>> &matrix) {
    int n = sz(matrix);
    vector<int> colOrder(n);
    vector<vector<double>>> inverse(n, vector<double>(n));
    for (int i = 0; i < (n); ++i) inverse[i][i] = 1, colOrder[i] = i;

    for (int i = 0; i < (n); ++i) {
        int pivotRow = i, pivotCol = i;
        for (int j = i; j < (n); ++j)
            for (int k = i; k < (n); ++k)
                if (fabs(matrix[j][k]) > fabs(matrix[pivotRow][pivotCol]))
                    pivotRow = j, pivotCol = k;
        if (fabs(matrix[pivotRow][pivotCol]) < 1e-12) return i;
        matrix[i].swap(matrix[pivotRow]);
        inverse[i].swap(inverse[pivotRow]);
        for (int j = 0; j < (n); ++j)
            swap(matrix[j][i], matrix[j][pivotCol]), swap(inverse[j][i], inverse[j][pivotCol]);
        swap(colOrder[i], colOrder[pivotCol]);
        double pivot = matrix[i][i];
        for (int j = i + 1; j < (n); ++j) {
            double factor = matrix[j][i] / pivot;
            matrix[j][i] = 0;
            for (int k = i + 1; k < (n); ++k) matrix[j][k] -= factor * matrix[i][k];
            for (int k = 0; k < (n); ++k) inverse[j][k] -= factor * inverse[i][k];
        }
    }
```

```
    }
    for (int j = i + 1; j < (n); ++j) matrix[i][j] /= pivot;
    for (int j = 0; j < (n); ++j) inverse[i][j] /= pivot;
    matrix[i][i] = 1;
}

for (int i = n - 1; i > 0; --i)
    for (int j = 0; j < (i); ++j) {
        double pivot = matrix[j][i];
        for (int k = 0; k < (n); ++k) inverse[j][k] -= pivot * inverse[i][k];
    }

for (int i = 0; i < (n); ++i)
    for (int j = 0; j < (n); ++j) matrix[colOrder[i]][colOrder[j]] = inverse[i][j];
return n;
}
```

Tridiagonal Linear System

Description: $x = \text{tridiagonal}(d, p, q, b)$ solves the equation system

$$\begin{pmatrix} b_0 \\ b_1 \\ b_2 \\ b_3 \\ \vdots \\ \vdots \\ b_{n-1} \end{pmatrix} = \begin{pmatrix} d_0 & p_0 & 0 & 0 & \cdots & 0 \\ q_0 & d_1 & p_1 & 0 & \cdots & 0 \\ 0 & q_1 & d_2 & p_2 & \cdots & 0 \\ \vdots & \vdots & \ddots & \ddots & \ddots & \vdots \\ 0 & 0 & \cdots & q_{n-3} & d_{n-2} & p_{n-2} \\ 0 & 0 & \cdots & 0 & q_{n-2} & d_{n-1} \end{pmatrix} \begin{pmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \\ \vdots \\ \vdots \\ x_{n-1} \end{pmatrix}.$$

This is useful for solving problems on the type

$$a_i = b_i a_{i-1} + c_i a_{i+1} + d_i, 1 \leq i \leq n,$$

where a_0, a_{n+1}, b_i, c_i and d_i are known. a can then be obtained from

$$\{a_i\} = \text{tridiagonal}(\{1, -1, -1, \dots, -1, 1\}, \{0, c_1, c_2, \dots, c_n\}, \{b_1, b_2, \dots, b_n, 0\}, \{a_0, d_1, d_2, \dots, d_n, a_{n+1}\}).$$

Fails if the solution is not unique.
If $|d_i| > |p_i| + |q_{i-1}|$ for all i , or $|d_i| > |p_{i-1}| + |q_i|$, or the matrix is positive definite, the algorithm is numerically stable and neither swapped nor the check for diag[i] == 0 is needed. Use when only the main and neighboring diagonals can be nonzero. diag and rhs have length n>=1; super/sub have length n-1. super[i] is A[i][i+1] and sub[i] is A[i+1][i]. Returns the solution, copying diag/rhs. Supply a nonsingular system and heed the stability conditions above.
Usage: auto x=tridiagonal({2,2},{-1},{-1},{1,1});
// approximately {1,1}
Time: $\mathcal{O}(N)$

```
typedef double T;
vector<T> tridiagonal(vector<T> diag, const vector<T> &super,
    const vector<T> &sub, vector<T> rhs) {
    int n = sz(rhs);
    vector<int> swapped(n);
    for (int i = 0; i < (n - 1); ++i) {
        if (abs(diag[i]) < 1e-9 * abs(super[i])) { // diag[i] == 0
            rhs[i + 1] -= rhs[i] * diag[i + 1] / super[i];
            if (i + 2 < n) rhs[i + 2] -= rhs[i] * sub[i + 1] / super[i + 1];
            diag[i + 1] = sub[i];
            swapped[++i] = 1;
        } else {
            diag[i + 1] -= super[i] * sub[i] / diag[i];
            rhs[i + 1] -= rhs[i] * sub[i] / diag[i];
        }
    }
    for (int i = n; i--;) {
        if (swapped[i]) {
            swap(rhs[i], rhs[i - 1]);
            diag[i - 1] = diag[i];
            rhs[i] /= super[i - 1];
        }
```

```
    } else {
        rhs[i] /= diag[i];
        if (i) rhs[i - 1] -= rhs[i] * super[i - 1];
    }
}
return rhs;
}
```

Number theory (4)

4.1 Modular arithmetic

Modular Arithmetic

Description: Set global mod first; residues are normalized to [0,mod) and read via .value. Division asserts unless its divisor is coprime with mod. Exponents must be nonnegative; multiplication must fit ll before reduction.

Usage: Mod a=5, b=3; ll residue=(a*b+a).value;
Mod power=a^10; Mod quotient=a/b; // gcd(b,mod)=1
"euclid.h"

```
const ll mod = 17; // change to something else
struct Mod {
    ll value;
    Mod(ll normalized) : Mod(normalized % mod + mod, 0) {}
    Mod(ll normalized, int) : value(normalized < mod ? normalized
        : normalized - mod) {}

    Mod operator+(Mod other) { return {value + other.value, 0}; }
    Mod operator-(Mod other) { return {value - other.value + mod,
        0}; }
    Mod operator*(Mod other) { return {value * other.value % mod,
        0}; }

    Mod operator/(Mod other) { return *this * invert(other); }
    Mod invert(Mod operand) {
        ll inverseCoeff, otherCoeff, gcdValue = euclid(operand.
            value, mod, inverseCoeff, otherCoeff);
        assert(gcdValue == 1);
        return inverseCoeff;
    }

    Mod operator^(ll exponent) {
        if (!exponent) return 1;
        Mod result = *this ^ (exponent / 2);
        result = result * result;
        return exponent & 1 ? *this * result : result;
    }
};
```

Modular Inverses

Description: Choose prime mod and LIM<=mod, then run the initialization loop inside setup/main. inverse[i] is valid for 1<=i<LIM; zero has no inverse. Costs $\mathcal{O}(\textit{LIM})$ time and memory.

Usage: // Define prime mod and LIM, then run the shown setup loop.

```
ll invOfThree=inverse[3]; // requires LIM>3
const ll mod = 1000000007, LIM = 200000;
ll *inverse = new ll[LIM] - 1;
inverse[1] = 1;
for (int i = 2; i < (LIM); ++i) inverse[i] = mod - (mod / i) *
    inverse[mod % i] % mod;
```

Modular Exponentiation

Description: Two-argument power modulo global mod. Use exponent>=0 and normalize base to [0,mod); products must fit ll. For prime mod and nonzero a, modpow(a,mod-2) is its inverse. For large moduli use the three-argument version in ModMulLL.h.

Usage: ll power=modpow(2,10); // 1024 modulo mod
ll inverse=modpow(3,mod-2); // requires prime mod, 3%mod!=0

```
const ll mod = 1000000007; // faster if const
ll modpow(ll base, ll exponent) {
    ll result = 1;
```

```
    for (; exponent; base = base * base % mod, exponent /= 2)
        if (exponent & 1) result = result * base % mod;
    return result;
}
```

Discrete Logarithm

Description: Returns the smallest $x > 0$ s.t. $a^x = b \pmod m$, or -1 if no such x exists. modLog(a,1,m) can be used to calculate the order of a . Use for discrete logarithms when sqrt(modulus) work/storage fit. Also supports non-coprime bases. The exponent is strictly positive: target=1 asks for a multiplicative order, not exponent zero. Normalize base and target to [0,modulus); modulus is positive and direct modular products must fit ll.

Usage: ll exponent=modLog(2,8,13); // 3
ll absent=modLog(2,3,4); // -1

```
Time:  $\mathcal{O}(\sqrt{m})$ 
ll modLog(ll base, ll target, ll modulus) {
    ll blockSize = (ll)sqrt(modulus) + 1, power = 1, giantStep =
        1, j = 1;
    unordered_map<ll, ll> babySteps;
    while (j <= blockSize && (power = giantStep = power * base %
        modulus) != target % modulus)
        babySteps[power * target % modulus] = j++;
    if (power == target % modulus) return j;
    if (__gcd(modulus, power) == __gcd(modulus, target))
        for (int i = 2; i < (blockSize + 2); ++i)
            if (babySteps.count(power = power * giantStep % modulus))
                return blockSize * i - babySteps[power];
    return -1;
}
```

Floor Sum and Modular Sum

Description: Computes progression sums without iterating: modsum(count,offset,slope,modulus) is $\sum_{i=0}^{count-1} ((offset + slope \cdot i) \bmod modulus)$ with nonnegative residues. divsum gives the corresponding sum of floored quotients for nonnegative inputs. Use for lattice-point counts or periodic sums. modulus must be positive; count is nonnegative. modsum normalizes negative offset/slope. Intermediate products must fit the implementation's integer arithmetic.

Usage: ll sum=modsum(4,1,2,5); // (1+3+0+2)=6
ll floors=divsum(4,1,2,5); // (0+0+1+1)=2

Time: $\log(m)$, with a large constant.

```
ull sumsq(ull count) { return count / 2 * ((count - 1) | 1); }

ull divsum(ull count, ull offset, ull slope, ull modulus) {
    ull answer = slope / modulus * sumsq(count) + offset /
        modulus * count;
    slope %= modulus;
    offset %= modulus;
    if (!slope) return answer;
    ull nextCount = (count * slope + offset) / modulus;
    return answer + (count - 1) * nextCount - divsum(nextCount,
        modulus - 1 - offset, modulus, slope);
}

ll modsum(ull count, ll offset, ll slope, ll modulus) {
    offset = ((offset % modulus) + modulus) % modulus;
    slope = ((slope % modulus) + modulus) % modulus;
    return count * offset + slope * sumsq(count) - modulus *
        divsum(count, offset, slope, modulus);
}
```

64-bit Modular Multiplication

Description: Calculate $a \cdot b$ mod c (or a^b mod c) for $0 \leq a, b \leq c \leq 7.2 \cdot 10^{18}$. Use for large-modulus exponentiation and primality/factorization when ordinary ll multiplication would overflow. The three-argument modpow takes an explicit modulus. The large bound assumes an 80-bit long double, typical with GNU on x86-64; double-width long double has the smaller bound explained in Details. Normalize operands first and use a positive modulus.

Usage: ull product=modmul(1000000000000ULL, 1000000000000ULL, 1000000000000003ULL);
ull power=modpow(2,100,1000000007);

Time: $\mathcal{O}(1)$ for modmul, $\mathcal{O}(\log b)$ for modpow

```
ull modmul(ull lhs, ull rhs, ull modulus) {
    ll remainder = lhs * rhs - modulus * (ull)(1.L / modulus *
        lhs * rhs);
    return remainder + modulus * (remainder < 0) - modulus * (
        remainder >= (ll)modulus);
}

ull modpow(ull base, ull exponent, ull mod) {
    ull result = 1;
    for (; exponent; base = modmul(base, base, mod), exponent /=
        2)
        if (exponent & 1) result = modmul(result, base, mod);
    return result;
}
```

Modular Square Root

Description: Tonelli-Shanks algorithm for modular square roots. Finds x s.t. $x^2 = a \pmod p$ ($-x$ gives the other solution). Use an odd prime p and a quadratic residue a ; test modpow(a,(p-1)/2,p)==1 for nonzero a first. Non-residues assert rather than return a failure marker. Handle $p=2$ separately. Requires the three-argument modpow from ModMulLL.h; all direct products must fit ll. The other root is $(p-x)\%p$.

Usage: ll x=sqrt(10LL,13LL); // 6 or 7

ll other=(13-x)%13;
Time: $\mathcal{O}(\log^2 p)$ worst case, $\mathcal{O}(\log p)$ for most p

```
"ModMulLL.h"
ll sqrt(ll a, ll prime) {
    a %= prime;
    if (a < 0) a += prime;
    if (a == 0) return 0;
    assert(modpow(a, (prime - 1) / 2, prime) == 1); // else no
        solution
    if (prime % 4 == 3) return modpow(a, (prime + 1) / 4, prime);
    // a^(n+3)/8 or 2^(n+3)/8 * 2^(n-1)/4 works if p % 8 == 5
    ll oddPart = prime - 1, nonResidue = 2;
    int twos = 0, order;
    while (oddPart % 2 == 0) ++twos, oddPart /= 2;
    while (modpow(nonResidue, (prime - 1) / 2, prime) != prime -
        1) ++nonResidue;
    ll x = modpow(a, (oddPart + 1) / 2, prime);
    ll remainder = modpow(a, oddPart, prime), rootOfUnity =
        modpow(nonResidue, oddPart, prime);
    for (; twos = order) {
        ll power = remainder;
        for (order = 0; order < twos && power != 1; ++order) power
            = power * power % prime;
        if (order == 0) return x;
        ll correction = modpow(rootOfUnity, 1LL << (twos - order -
            1), prime);
        rootOfUnity = correction * correction % prime;
        x = x * correction % prime;
        remainder = remainder * rootOfUnity % prime;
    }
}
```

4.2 Primality

Linear Sieve (Smallest Prime Factors)

Description: lp[x] is the smallest prime factor for $2 \leq x \leq n$; lp[0]=lp[1]=0. Each composite is generated once. Uses more memory than the boolean sieve, but supports fast factorization.

Usage: LinearSieve sieve(100);
int smallest=sieve.lp[84]; // 2
int x=84;
while (x>1) { cout << sieve.lp[x] << ' '; x/=sieve.lp[x]; }
// 2 2 3 7

Time: $\mathcal{O}(N)$ construction and memory; $\mathcal{O}(\log x)$ to factor x using lp.

```
struct LinearSieve {
    vector<int> lp, primes;
    LinearSieve(int n) : lp(n + 1) {
        for (int i = 2; i <= n; i++) {
            if (lp[i] == 0) lp[i] = i, primes.push_back(i);
            for (int p : primes) {
                if (1LL * i * p > n) break;
                lp[i * p] = p;
                if (p == lp[i]) break;
            }
        }
    }
};
```

Fast Prime Sieve

Description: Prime sieve for generating all primes smaller than LIM. Set LIM before compilation and call eratosthenes once. Returns primes in increasing order and fills the global isPrime bitset below LIM. This segmented version skips even candidates to reduce memory traffic; use when generating a large prime list.

Usage: auto primes=eratosthenes();
bool prime=isPrime[97]; // requires LIM>97
Time: LIM=1e9 $\approx$ 1.5s

```
const int LIM = 1e6;
bitset<LIM> isPrime;
vector<int> eratosthenes() {
    const int blockSize = (int)round(sqrt(LIM)), oddCount = LIM / 2;
    vector<int> primes = {2}, sieve(blockSize + 1);
    primes.reserve((int)(LIM / log(LIM) * 1.1));
    vector<pii> sievingPrimes;
    for (int i = 3; i <= blockSize; i += 2)
        if (!sieve[i]) {
            sievingPrimes.push_back({i, i * i / 2});
            for (int j = i * i; j <= blockSize; j += 2 * i) sieve[j] = 1;
        }
    for (int blockStart = 1; blockStart <= oddCount; blockStart += blockSize) {
        array<bool, blockSize> block{};
        for (auto &p, nextMultiple : sievingPrimes)
            for (int i = nextMultiple; i < blockStart + blockSize;
                nextMultiple = (i += p))
                block[i - blockStart] = 1;
        for (int i = 0; i < (min(blockSize, oddCount - blockStart))
            ; ++i)
            if (!block[i]) primes.push_back((blockStart + i) * 2 + 1);
    }
    for (int i : primes) isPrime[i] = 1;
    return primes;
}
```

Miller-Rabin Primality Test

Description: Deterministic Miller-Rabin primality test. Guaranteed to work for numbers up to $7 \cdot 10^{18}$; for larger numbers, use Python and extend A randomly. Use when testing individual large integers without sieving to them. Returns false for 0 and 1. A fixed witness set gives the stated deterministic range, using ModMulLL for modular products. The name isPrime conflicts with the sieve bitset if both snippets are pasted unchanged.
Usage: bool prime=isPrime(1000000007ULL); // true
Time: 7 times the complexity of $a^b \bmod c$.

```
bool isPrime(ull n) {
    if (n < 2 || n % 6 % 4 != 1) return (n | 1) == 3;
    ull bases[] = {2, 325, 9375, 28178, 450775, 9780504,
        1795265022}, twos = __builtin_ctzll(n - 1),
```

```
    oddPart = n >> twos;
    for (ull base : bases) { // ^ count trailing zeroes
        ull power = modpow(base % n, oddPart, n), squaresLeft = twos;
        while (power != 1 && power != n - 1 && base % n && squaresLeft--)
            power = modmul(power, power, n);
        if (power != n - 1 && squaresLeft != twos) return 0;
    }
    return 1;
}
```

Pollard-Rho Factorization

Description: Pollard-rho randomized factorization algorithm. Returns prime factors of a number, in arbitrary order (e.g. 2299 -> {11, 19, 11}). Call factor on a positive integer; factor(1) is empty and factor(0) is unsupported. Repeated factors are retained and output is unsorted. Sort to group prime powers for divisor counts or totients. pollard is an internal splitter for composite n, not a primality test. ModMulLL platform/range limits apply.
Usage: auto factors=factor(2299); sort(all(factors));
// {11,11,19}
Time: $\mathcal{O}\left(n^{1/4}\right)$, less for numbers with small factors.
"ModMulLL.h", "MillerRabin.h"

```
ull pollard(ull n) {
    ull slow = 0, fast = 0, iterations = 30, product = 2, seed = 1, nextProduct;
    auto advance = [&](ull slow) { return modmul(slow, slow, n) + seed; };
    while (iterations++ % 40 || __gcd(product, n) == 1) {
        if (slow == fast) slow = ++seed, fast = advance(slow);
        if ((nextProduct = modmul(product, max(slow, fast) - min(slow, fast), n))
            product = nextProduct;
            slow = advance(slow), fast = advance(advance(fast)));
    }
    return __gcd(product, n);
}
vector<ull> factor(ull n) {
    if (n == 1) return {};
    if (isPrime(n)) return {n};
    ull divisor = pollard(n);
    auto leftFactors = factor(divisor), rightFactors = factor(n / divisor);
    leftFactors.insert(leftFactors.end(), all(rightFactors));
    return leftFactors;
}
```

4.3 Divisibility

Extended Euclidean Algorithm

Description: Returns $g = \gcd(a, b)$ and coefficients with $ax + by = g$. Use positive a,b. For $ax+by=c$ require g to divide c, then scale coefficients by c/g. For an inverse modulo m require g=1 and normalize x to [0,m). Products must fit ll.
Usage: ll x,y; ll q=euclid(30,18,x,y); // 30*x+18*y=6

```
ll euclid(ll a, ll b, ll &x, ll &y) {
    if (!b) return x = 1, y = 0, a;
    ll gcdValue = euclid(b, a % b, y, x);
    return y -= a / b * x, gcdValue;
}
```

Chinese Remainder Theorem (CRT)

Description: crt(a,m,b,n) returns x with $x \equiv a \pmod m$, $x \equiv b \pmod n$. Use positive moduli and normalized residues; require $(a-b) \bmod \gcd(m,n) = 0$ or it asserts. Returns $0 \leq x < \text{lcm}(m,n)$, including non-coprime moduli. Requires $mn < 2^{62}$.
Usage: ll x=crt(2,3,3,5); // 8 modulo 15
// Check (a-b)%gcd(m,n)==0 before calling.

```
Time: log(n)
"euclid.h"
ll crt(ll a, ll m, ll b, ll n) {
    if (n > m) swap(a, b), swap(m, n);
    ll x, y, gcdValue = euclid(m, n, x, y);
    assert((a - b) % gcdValue == 0); // else no solution
    x = (b - a) % n * x % n / gcdValue * m + a;
    return x < 0 ? x + m * n / gcdValue : x;
}
```

4.3.1 Bézout’s identity

For $a \neq, b \neq 0$, then $d = \gcd(a, b)$ is the smallest positive integer for which there are integer solutions to

$$ax + by = d$$

If (x, y) is one solution, then all solutions are given by

$$\left(x + \frac{kb}{\gcd(a,b)}, y - \frac{ka}{\gcd(a,b)}\right), \quad k \in \mathbb{Z}$$

Euler’s Totient Function

Description: Choose LIM and call calculatePhi once for phi[1..LIM-1]. phi[n] counts residues coprime with n; $\phi(n) = n \prod_{p|n} (1 - 1/p)$. Reduce exponents modulo phi(n) directly only when base and n are coprime.
Usage: calculatePhi(); int count=phi[12]; // 4

```
const int LIM = 5000000;
int phi[LIM];
void calculatePhi() {
    for (int i = 0; i < (LIM); ++i) phi[i] = i & 1 ? i : i / 2;
    for (int i = 3; i < LIM; i += 2)
        if (phi[i] == i)
            for (int j = i; j < LIM; j += i) phi[j] -= phi[j] / i;
}
```

4.4 Fractions

Continued Fractions

Description: Given N and a real number $x \geq 0$, finds the closest rational approximation p/q with $p, q \leq N$. It will obey $|p/q - x| \leq 1/qN$. For consecutive convergents, $p_{k+1}q_k - q_{k+1}p_k = (-1)^k$. (p_k/q_k alternates between $> x$ and $< x$.) If x is rational, y eventually becomes ∞ ; if x is the root of a degree 2 polynomial the a ’s eventually become cyclic. Use to approximate nonnegative real input by a bounded rational. limit $\geq$ 1 bounds both numerator and denominator, not only the denominator. Returns (p,q). Precision comes from type d; use long double when increasing the limit substantially.
Usage: auto [p,q]=approximate(3.141592653589793,100);
// {22,7}
Time: $\mathcal{O}(\log N)$

```
typedef double d; // for N ~ 1e7; long double for N ~ 1e9
pair<ll, ll> approximate(d x, ll limit) {
    ll prevNum = 0, prevDen = 1, numerator = 1, denominator = 0,
        inf = LLONG_MAX;
    d remainder = x;
    for (;) {
        ll maxStep = min(numerator ? (limit - prevNum) / numerator
            : inf,
            denominator ? (limit - prevDen) /
                denominator : inf),
            quotient = (ll)floor(remainder), step = min(quotient,
                maxStep),
            nextNum = step * numerator + prevNum, nextDen = step *
                denominator + prevDen;
        if (quotient > step) {
```

```
// If step > quotient/2, we have a semi-convergent that
// gives us a
// better approximation; if step = quotient/2, we *may*
// have one.
// Return {numerator, denominator} here for a more
// canonical approximation.
return (abs(x - (d)nextNum / (d)nextDen) < abs(x - (d)
numerator / (d)denominator))
    ? make_pair(nextNum, nextDen)
    : make_pair(numerator, denominator);
}
if (abs(remainder = 1 / (remainder - (d)quotient)) > 3 *
limit) return {nextNum, nextDen};
prevNum = numerator;
numerator = nextNum;
prevDen = denominator;
denominator = nextDen;
}
}
```

Fraction Binary Search

Description: Given f and N , finds the smallest fraction $p/q \in [0, 1]$ such that $f(p/q)$ is true, and $p, q \leq N$. You may want to throw an exception from f if it finds an exact solution, in which case N can be removed. The predicate must be monotone false-to-true as p/q increases, and true at 1. limit>=1 bounds both p and q. Compare fractions by cross multiplication to avoid rounding; products must fit ll. Finds an exact bounded fraction rather than a floating approximation.
Usage: Frac answer=fracBS([](Frac x) { return 3*x.p>=x.q; },10);
// answer is {1,3}
Time: $\mathcal{O}(\log(N))$

```
struct Frac {
    ll p, q;
};
template <class F> Frac fracBS(F predicate, ll limit) {
    bool moveHi = 1, prevAdvanced = 1, advanced = 1;
    Frac lo{0, 1}, hi{1, 1}; // Set hi to 1/0 to search (0, N]
    if (predicate(lo)) return lo;
    assert(predicate(hi));
    while (prevAdvanced || advanced) {
        ll advance = 0, step = 1; // move hi if moveHi, else lo
        for (int shift = 0; step; (step *= 2) >= shift) {
            advance += step;
            Frac mid{lo.p * advance + hi.p, lo.q * advance + hi.q};
            if (abs(mid.p) > limit || mid.q > limit || moveHi == !
                predicate(mid)) {
                advance -= step;
                shift = 2;
            }
        }
        hi.p += lo.p * advance;
        hi.q += lo.q * advance;
        moveHi = !moveHi;
        swap(lo, hi);
        prevAdvanced = advanced;
        advanced = !advance;
    }
    return moveHi ? hi : lo;
}
```

4.5 Pythagorean Triples

The Pythagorean triples are uniquely generated by

$$a = k \cdot (m^2 - n^2), \quad b = k \cdot (2mn), \quad c = k \cdot (m^2 + n^2),$$

with $m > n > 0, k > 0, m \perp n$, and either m or n even.

4.6 Primes

$p = 962592769$ is such that $2^{21} \mid p - 1$, which may be useful. For hashing use 970592641 (31-bit number), 31443539979727 (45-bit), 3006703054056749 (52-bit). There are 78498 primes less than 1 000 000.

Primitive roots exist modulo any prime power p^a , except for $p = 2, a > 2$, and there are $\phi(\phi(p^a))$ many. For $p = 2, a > 2$, the group $\mathbb{Z}_{2^a}^\times$ is instead isomorphic to $\mathbb{Z}_2 \times \mathbb{Z}_{2^{a-2}}$.

4.7 Estimates

$$\sum_{d \mid n} d = O(n \log \log n).$$

The number of divisors of n is at most around 100 for $n < 5e4$, 500 for $n < 1e7$, 2000 for $n < 1e10$, 200 000 for $n < 1e19$.

4.8 Mobius Function

$$\mu(n) = \begin{cases} 0 & n \text{ is not square free} \\ 1 & n \text{ has even number of prime factors} \\ -1 & n \text{ has odd number of prime factors} \end{cases}$$

Mobius Inversion:

$$g(n) = \sum_{d \mid n} f(d) \Leftrightarrow f(n) = \sum_{d \mid n} \mu(d) g(n/d)$$

Other useful formulas/forms:

$$\sum_{d \mid n} \mu(d) = [n = 1] \text{ (very useful)}$$

$$g(n) = \sum_{m \mid d} f(d) \Leftrightarrow f(n) = \sum_{n \mid d} \mu(d/n) g(d)$$

$$g(n) = \sum_{1 \leq m \leq n} f(\lfloor \frac{n}{m} \rfloor) \Leftrightarrow f(n) = \sum_{1 \leq m \leq n} \mu(m) g(\lfloor \frac{n}{m} \rfloor)$$

Combinatorial (5)

5.1 Permutations

5.1.1 Factorial

n	1	2	3	4	5	6	7	8	9	10
$n!$	1	2	6	24	120	720	5040	40320	362880	3628800
n	11	12	13	14	15	16	17			
$n!$	4.0e7	4.8e8	6.2e9	8.7e10	1.3e12	2.1e13	3.6e14			
n	20	25	30	40	50	100	150	171		
$n!$	2e18	2e25	3e32	8e47	3e64	9e157	6e262	>DBL_MAX		

Permutation Ranking

Description: Maps a permutation of 0..n-1 to a unique ID in [0,n!), without preserving lexicographic order. The return type is int; use n<=12.
Usage: vector<int> p={2,0,1}; int state=permToInt(p);
Time: $\mathcal{O}(n)$

```
int permToInt(vector<int> &permutation) {
    int usedMask = 0, i = 0, rank = 0;
    for (int value : permutation)
        rank = rank * ++i + __builtin_popcount(usedMask & ~(1 <<
            value)),
            usedMask |= 1 << value; // (note: minus, not ~!)
    return rank;
}
```

5.1.2 Cycles

Let $g_S(n)$ be the number of n -permutations whose cycle lengths all belong to the set S . Then

$$\sum_{n=0}^\infty g_S(n) \frac{x^n}{n!} = \exp \left(\sum_{n \in S} \frac{x^n}{n} \right)$$

5.1.3 Derangements

Permutations of a set such that none of the elements appear in their original position.

$$D(n) = (n-1)(D(n-1) + D(n-2)) = nD(n-1) + (-1)^n = \left\lfloor \frac{n!}{e} \right\rfloor$$

5.1.4 Burnside’s lemma

Given a group G of symmetries and a set X , the number of elements of X up to symmetry equals

$$\frac{1}{|G|} \sum_{g \in G} |X^g|,$$

where X^g are the elements fixed by g ($g.x = x$).

If $f(n)$ counts “configurations” (of some sort) of length n , we can ignore rotational symmetry using $G = \mathbb{Z}_n$ to get

$$g(n) = \frac{1}{n} \sum_{k=0}^{n-1} f(\gcd(n, k)) = \frac{1}{n} \sum_{k \mid n} f(k) \phi(n/k).$$

5.2 Partitions and subsets

5.2.1 Partition function

Number of ways of writing n as a sum of positive integers, disregarding the order of the summands.

$$p(0) = 1, \quad p(n) = \sum_{k \in \mathbb{Z} \setminus \{0\}} (-1)^{k+1} p(n - k(3k - 1)/2)$$

$$p(n) \sim 0.145/n \cdot \exp(2.56\sqrt{n})$$

n	0	1	2	3	4	5	6	7	8	9	20	50	100
$p(n)$	1	1	2	3	5	7	11	15	22	30	627	~2e5	~2e8

5.2.2 Lucas’ Theorem

Let n, m be non-negative integers and p a prime. Write $n = n_k p^k + \dots + n_1 p + n_0$ and $m = m_k p^k + \dots + m_1 p + m_0$. Then $\binom{n}{m} \equiv \prod_{i=0}^k \binom{n_i}{m_i} \pmod{p}$.

5.2.3 Binomials

Multinomial Coefficients

Description: Counts multiset arrangements: $\frac{(\sum k_i)!}{\prod k_i!}$. counts[i] must be nonnegative. Exact integer calculation without a modulus; the result and intermediate products must fit ll.

Usage: vector<int> counts={2,1}; ll ways=multinomial(counts);
// 3

```
ll multinomial(vector<int> &counts) {
```

```
ll ways = 1, total = counts.empty() ? 1 : counts[0];
for (int i = 1; i < (sz(counts)); ++i)
    for (int j = 0; j < (counts[i]); ++j) ways = ways * ++total
        / (j + 1);
return ways;
}
```

5.3 General purpose numbers

5.3.1 Bernoulli numbers

EGF of Bernoulli numbers is $B(t) = \frac{t}{e^t - 1}$ (FFT-able).
 $B[0, \dots] = [1, -\frac{1}{2}, \frac{1}{6}, 0, -\frac{1}{30}, 0, \frac{1}{42}, \dots]$

Sums of powers:

$$\sum_{i=1}^n n^m = \frac{1}{m+1} \sum_{k=0}^m \binom{m+1}{k} B_k \cdot (n+1)^{m+1-k}$$

Euler-Maclaurin formula for infinite sums:

$$\sum_{i=m}^\infty f(i) = \int_m^\infty f(x)dx - \sum_{k=1}^\infty \frac{B_k}{k!} f^{(k-1)}(m) \\ \approx \int_m^\infty f(x)dx + \frac{f(m)}{2} - \frac{f'(m)}{12} + \frac{f'''(m)}{720} + O(f^{(5)}(m))$$

5.3.2 Stirling numbers of the first kind

Number of permutations on n items with k cycles.

$$c(n, k) = c(n-1, k-1) + (n-1)c(n-1, k), \quad c(0, 0) = 1 \\ \sum_{k=0}^n c(n, k) x^k = x(x+1) \dots (x+n-1)$$

$c(8, k) = 8, 0, 5040, 13068, 13132, 6769, 1960, 322, 28, 1$
 $c(n, 2) = 0, 0, 1, 3, 11, 50, 274, 1764, 13068, 109584, \dots$

5.3.3 Eulerian numbers

Number of permutations $\pi \in S_n$ in which exactly k elements are greater than the previous element. k j :s s.t. $\pi(j) > \pi(j+1)$, $k+1$ j :s s.t. $\pi(j) \geq j$, k j :s s.t. $\pi(j) > j$.

$$E(n, k) = (n-k)E(n-1, k-1) + (k+1)E(n-1, k)$$

$$E(n, 0) = E(n, n-1) = 1$$

$$E(n, k) = \sum_{j=0}^k (-1)^j \binom{n+1}{j} (k+1-j)^n$$

5.3.4 Stirling numbers of the second kind

Partitions of n distinct elements into exactly k groups.

$$S(n, k) = S(n-1, k-1) + kS(n-1, k)$$

$$S(n, 1) = S(n, n) = 1$$

$$S(n, k) = \frac{1}{k!} \sum_{j=0}^k (-1)^{k-j} \binom{k}{j} j^n$$

5.3.5 Bell numbers

Total number of partitions of n distinct elements. $B(n) = 1, 1, 2, 5, 15, 52, 203, 877, 4140, 21147, \dots$ For p prime,

$$B(p^m + n) \equiv mB(n) + B(n+1) \pmod{p}$$

5.3.6 Labeled unrooted trees

on n vertices: n^{n-2}
on k existing trees of size n_i : $n_1 n_2 \dots n_k n^{k-2}$
with degrees d_i : $(n-2)! / ((d_1-1)! \dots (d_n-1)!)$

5.3.7 Catalan numbers

$$C_n = \frac{1}{n+1} \binom{2n}{n} = \binom{2n}{n} - \binom{2n}{n+1} = \frac{(2n)!}{(n+1)!n!}$$

$$C_0 = 1, \quad C_{n+1} = \frac{2(2n+1)}{n+2} C_n, \quad C_{n+1} = \sum C_i C_{n-i}$$

$C_n = 1, 1, 2, 5, 14, 42, 132, 429, 1430, 4862, 16796, 58786, \dots$

- sub-diagonal monotone paths in an $n \times n$ grid.
- strings with n pairs of parenthesis, correctly nested.
- binary trees with with $n+1$ leaves (0 or 2 children).
- ordered trees with $n+1$ vertices.
- ways a convex polygon with $n+2$ sides can be cut into triangles by connecting vertices with straight lines.
- permutations of $[n]$ with no 3-term increasing subseq.

Data structures (6)

Ordered Set (PBDS)

Description: GNU PBDS set with 0-based ranks. `order_of_key(x)` counts elements $<x$; `find_by_order(k)` returns an iterator, or end if absent. For duplicates use `(value,uniqueId)`, not `less.equal`. Replace `null_type` for a map.
Usage: `OrderedSet<int> s; s.insert(8); s.insert(10);`
`int smaller=s.order_of_key(10); // 1`
`int kth=s.find_by_order(0); // 8`

```
#include <bits/extc++.h>
using namespace __gnu_pbds;

template <class T>
using OrderedSet = tree<T, null_type, less<T>, rb_tree_tag,
    tree_order_statistics_node_update>;
void example() {
    OrderedSet<int> ordered, other;
    ordered.insert(8);
    auto inserted = ordered.insert(10).first;
    assert(inserted == ordered.lower_bound(9));
    assert(ordered.order_of_key(10) == 1);
    assert(ordered.order_of_key(11) == 2);
    assert(*ordered.find_by_order(0) == 8);
    ordered.join(other); // assuming T< T2 or T> T2, merge t2
                          into t
}
```

Custom Hash Map

Description: Integer-key map with randomized SplitMix64: fixed collision attacks are harder, but worst-case time is not guaranteed. Collisions affect speed, not correctness. CustomHash also works with `unordered_map`; only the displayed container needs GNU PBDS. Values are `int`. `hashMap[key]` inserts zero if absent; `find` does not insert. Iteration is unsorted.

Usage: `hashMap[42]++;`
`auto it=hashMap.find(42);`
`if (it!=hashMap.end()) cout << it->second;`
`hashMap.clear();` // between test cases
// Standard-library alternative:
`unordered_map<ll,int,CustomHash> freq;`
`freq.reserve(2*n);` // optional, reduces rehashing
Time: Expected $\mathcal{O}(1)$ per insert/find/erase; worst-case $\mathcal{O}(N)$ per operation.

```
#include <ext/pb_ds/assoc_container.hpp>
struct CustomHash {
    static ull mix(ull x) {
        // Fixed constants mix all 64 bits; unsigned overflow is
        // intentional.
        x += 0x9e3779b97f4a7c15ULL;
        x = (x ^ (x >> 30)) * 0xbf58476d1ce4e5b9ULL;
        x = (x ^ (x >> 27)) * 0x94d049bb133111ebULL;
        return x ^ (x >> 31);
    }
    size_t operator()(ll key) const {
        static const ull seed = chrono::steady_clock::now().
            time_since_epoch().count();
        return mix((ull)key + seed);
    }
};
__gnu_pbds::gp_hash_table<ll, int, CustomHash> hashMap;
```

Iterative Segment Tree

Description: KACTL's 2N iterative maximum tree: 1-based point assignment and inclusive `[l,r]`. Initially uniform `initialValue`; no power-of-two padding. Empty queries return unit. For another aggregate, edit `T/merge/unit`; preserve the two accumulator orders. Initialize with unit then update leaves, or rebuild parents if the aggregate changes.

Usage: `IterativeSegTree seg(5,0);`
`seg.update(3,7);` // `a[3]=7`, not `+=7`
`int best=seg.query(2,4);` // 7; includes positions 2,3,4
`int single=seg.query(3,3);` // 7
// To build from `a[1..n]`: `update(i,a[i])` for `i=1..n`.
Time: $\mathcal{O}(N)$ construction, $\mathcal{O}(\log N)$ per update/query; $\mathcal{O}(N)$ memory (2N values).

```
struct IterativeSegTree {
    typedef int T;
    static constexpr T unit = INT_MIN;
    T merge(T a, T b) const { return max(a, b); }
    vector<T> seg;
    int n;
    IterativeSegTree(int n = 0, T initialValue = unit) : seg(2 *
        n, initialValue), n(n) {}
    void update(int pos, T val) {
        for (seg[pos += n - 1] = val; pos /= 2;) seg[pos] = merge(
            seg[pos * 2], seg[pos * 2 + 1]);
    }
    T query(int l, int r) const {
        if (n == 0 || l > r) return unit;
        T leftResult = unit, rightResult = unit;
        for (l += n - 1, r += n - 1; l <= r; l /= 2, r /= 2) {
            if (l & 1) leftResult = merge(leftResult, seg[l++]);
            if (!(r & 1)) rightResult = merge(seg[r--], rightResult);
        }
        return merge(leftResult, rightResult);
    }
};
```

Segment Tree + Walk

Description: Point assignment and range maximum. Positions are 1-based; query(l,r) is inclusive. Initially every position equals initialValue. To change the aggregate, edit T, merge and its identity unit; use a wider T if needed. firstAtLeast(l,r,val) walks left first to find the smallest position in [l,r] with value>=val, or -1 if absent. A node with maximum<val is skipped. This search relies on the maximum aggregate; changing merge also requires changing the walk.

Usage: SegTree seg(5,0); seg.update(3,7);
int best=seg.query(2,4); // 7; positions 2,3,4
int pos=seg.firstAtLeast(2,5,6); // 3; first value >=6 in [2,5]
// Find the first available capacity >=need, then consume it:
int need=4; pos=seg.firstAtLeast(1,5,need);
if (pos!=-1) seg.update(pos,seg.query(pos,pos)-need);
Time: $\mathcal{O}(\log N)$ per operation; $\mathcal{O}(N)$ memory.

```
struct SegTree {
    typedef int T;
    static constexpr T unit = INT_MIN;
    int n;
    vector<T> seg;
    SegTree(int n = 0, T initialValue = unit) : n(n), seg(4 * n + 4, unit) {
        if (n > 0) build(1, 1, n, initialValue);
    }
    T merge(T a, T b) const { return max(a, b); }
    void build(int id, int lo, int hi, T val) {
        if (lo == hi) {
            seg[id] = val;
            return;
        }
        int mid = (lo + hi) / 2;
        build(id * 2, lo, mid, val);
        build(id * 2 + 1, mid + 1, hi, val);
        seg[id] = merge(seg[id * 2], seg[id * 2 + 1]);
    }
    void update(int id, int lo, int hi, int pos, T val) {
        if (lo == hi) {
            seg[id] = val;
            return;
        }
        int mid = (lo + hi) / 2;
        if (pos <= mid)
            update(id * 2, lo, mid, pos, val);
        else
            update(id * 2 + 1, mid + 1, hi, pos, val);
        seg[id] = merge(seg[id * 2], seg[id * 2 + 1]);
    }
    void update(int pos, T val) { update(1, 1, n, pos, val); }
    T query(int id, int lo, int hi, int l, int r) const {
        if (r < lo || hi < l) return unit;
        if (l <= lo && hi <= r) return seg[id];
        int mid = (lo + hi) / 2;
        return merge(query(id * 2, lo, mid, l, r), query(id * 2 + 1, mid + 1, hi, l, r));
    }
    T query(int l, int r) const {
        if (n == 0 || l > r) return unit;
        return query(1, 1, n, l, r);
    }
    int firstAtLeast(int id, int lo, int hi, int l, int r, T val) const {
        if (r < lo || hi < l || seg[id] < val) return -1; // No valid position here.
        if (lo == hi) return lo;
        int mid = (lo + hi) / 2;
        int pos = firstAtLeast(id * 2, lo, mid, l, r, val);
        if (pos != -1) return pos; // Left first gives the smallest position.
    }
};
```

Lazy Segment Tree

Description: Range assignment, addition and maximum on inclusive, 1-based [l,r]. The vector constructor ignores values[0]; default initial value is -inf, so initialize finite values before addition. Values and additions must fit int. Keep hasSet for assignments; changing to sums also requires segment lengths.

Usage: LazySegTree seg(5,0); seg.add(1,3,4); seg.set(2,2,9);
int best=seg.query(1,5); // 9
Time: $\mathcal{O}(\log N)$ per operation; $\mathcal{O}(N)$ memory.

```
const int inf = 1000000000;
struct LazySegTree {
    struct Node {
        int val = -inf;
        int lazySet = 0;
        int lazyAdd = 0;
        bool hasSet = false;
    };
    int n;
    vector<Node> seg;
    LazySegTree(int n, int initialValue = -inf) : n(n), seg(4 * n + 4, Node(initialValue)) {}
    LazySegTree(const vector<int> &values) : LazySegTree(max(0, sz(values) - 1)) {
        if (n > 0) build(1, 1, n, values);
    }
    void build(int id, int lo, int hi, const vector<int> &values) {
        if (lo == hi) {
            seg[id].val = values[lo];
            return;
        }
        int mid = (lo + hi) / 2;
        build(id * 2, lo, mid, values);
        build(id * 2 + 1, mid + 1, hi, values);
        pull(id);
    }
    void pull(int id) { seg[id].val = max(seg[id * 2].val, seg[id * 2 + 1].val); }
    void applySet(int id, int val) {
        seg[id].val = val;
        seg[id].lazySet = val;
        seg[id].lazyAdd = 0;
        seg[id].hasSet = true;
    }
    void applyAdd(int id, int val) {
        seg[id].val += val;
        if (seg[id].hasSet)
            seg[id].lazySet += val;
        else
            seg[id].lazyAdd += val;
    }
    void push(int id) {
        // Assignment overwrites older updates; addition comes after assignment.
        if (seg[id].hasSet) {
            applySet(id * 2, seg[id].lazySet);
            applySet(id * 2 + 1, seg[id].lazySet);
            seg[id].hasSet = false;
        }
        if (seg[id].lazyAdd != 0) {
            applyAdd(id * 2, seg[id].lazyAdd);
        }
    }
};
```

```
        applyAdd(id * 2 + 1, seg[id].lazyAdd);
        seg[id].lazyAdd = 0;
    }
}
void update(int id, int lo, int hi, int l, int r, int val, bool isAssignment) {
    if (r < lo || hi < l) return;
    if (l <= lo && hi <= r) {
        if (isAssignment)
            applySet(id, val);
        else
            applyAdd(id, val);
        return;
    }
    push(id);
    int mid = (lo + hi) / 2;
    update(id * 2, lo, mid, l, r, val, isAssignment);
    update(id * 2 + 1, mid + 1, hi, l, r, val, isAssignment);
    pull(id);
}
void set(int l, int r, int val) {
    if (n > 0 && l <= r) update(1, 1, n, l, r, val, true);
}
void add(int l, int r, int val) {
    if (n > 0 && l <= r) update(1, 1, n, l, r, val, false);
}
int query(int id, int lo, int hi, int l, int r) {
    if (r < lo || hi < l) return INT_MIN;
    if (l <= lo && hi <= r) return seg[id].val;
    push(id);
    int mid = (lo + hi) / 2;
    return max(query(id * 2, lo, mid, l, r), query(id * 2 + 1, mid + 1, hi, l, r));
}
int query(int l, int r) {
    if (n == 0 || l > r) return INT_MIN;
    return query(1, 1, n, l, r);
}
};
```

Persistent Segment Tree

Description: Persistent point assignment and range maximum on 1..N, inclusive [l,r]. Version 0 is uniformly initialValue (default 0). update(version,pos,val) copies the changed root-to-leaf path, returns a new version ID, and preserves every old version; branching is allowed. Node 0 represents an unchanged uniform subtree. Use node IDs, not references across push_back. Change max and unit for another aggregate; sums also need segment lengths for uniform subtrees.

Usage: PersistentSegTree seg(n); // initially all zero; require n>=1
int v1=seg.update(0,3,7); // assign a[3]=7
int v2=seg.update(v1,2,9);
int branch=seg.update(0,1,-5); // branch from the original version
ll best=seg.query(v2,2,3); // 9; seg.query(v1,2,3) is still 7
ll old=seg.query(0,1,n); // 0; positions and bounds are 1-based
Time: $\mathcal{O}(\log N)$ per update/query; $\mathcal{O}(1 + U \log N)$ memory for U updates.

```
struct PersistentSegTree {
    static constexpr ll unit = LLONG_MIN;
    struct Node {
        int left, right;
        ll val;
    };
    int n;
    vector<Node> nodes;
    vector<int> roots = {0};
    PersistentSegTree(int n, ll initialValue = 0) : n(n), nodes(1, {0, 0, initialValue}) {}
    int updateNode(int id, int lo, int hi, int pos, ll val) {
```

```
int cur = sz(nodes);
nodes.push_back(nodes[id]);
if (lo == hi) {
    nodes[cur].val = val;
    return cur;
}
int mid = (lo + hi) / 2;
if (pos <= mid)
    nodes[cur].left = updateNode(nodes[id].left, lo, mid, pos, val);
else
    nodes[cur].right = updateNode(nodes[id].right, mid + 1, hi, pos, val);
nodes[cur].val = max(nodes[nodes[cur].left].val, nodes[nodes[cur].right].val);
return cur;
}
int update(int version, int pos, ll val) {
    roots.push_back(updateNode(roots[version], 1, n, pos, val));
    return sz(roots) - 1;
}
ll queryNode(int id, int lo, int hi, int l, int r) const {
    if (r < lo || hi < l) return unit;
    if (id == 0 || (l <= lo && hi <= r)) return nodes[id].val;
    int mid = (lo + hi) / 2;
    return max(queryNode(nodes[id].left, lo, mid, l, r), queryNode(nodes[id].right, mid + 1, hi, l, r));
}
ll query(int version, int l, int r) const {
    if (l > r) return unit;
    return queryNode(roots[version], 1, n, l, r);
}
};
```

Disjoint Set Union (DSU)

Description: Union by size with path compression, for vertices 1..n. get returns the component representative; s[get(u)] is its size. This version cannot roll back merges.
Usage: DSU dsu(5);
dsu.join(1,2);
bool same=dsu.joined(1,2); // true
int size=dsu.s[dsu.get(1)]; // 2
Time: $\mathcal{O}(N)$ initialization; amortized $\mathcal{O}(\alpha(N))$ per operation.

```
struct DSU {
    vector<int> par, s;
    DSU(int n) { init(n); }
    void init(int n) {
        par.resize(n + 1);
        iota(all(par), 0);
        s.assign(n + 1, 1);
    }
    int get(int u) { return par[u] == u ? u : par[u] = get(par[u]); }
    void join(int u, int v) {
        u = get(u), v = get(v);
        if (u == v) return;
        if (s[u] < s[v]) swap(u, v);
        par[v] = u;
        s[u] += s[v];
    }
    bool joined(int u, int v) { return get(u) == get(v); }
};
```

Rollback DSU

Description: Rollback DSU on 0..n-1; allocate n+1 for labels 1..n. Save time() before temporary joins and rollback(checkpoint) afterward. Checkpoints are history lengths, not join counts. Do not add path compression unless its changes are also recorded.
Usage: RollbackDSU uf(n); int checkpoint=uf.time();
uf.join(u,v); bool connected=uf.find(u)==uf.find(v);
uf.rollback(checkpoint);
Time: $\mathcal{O}(\log(N))$

```
struct RollbackDSU {
    vector<int> parentOrSize;
    vector<pii> history;
    RollbackDSU(int n) : parentOrSize(n, -1) {}
    int size(int x) { return -parentOrSize[find(x)]; }
    int find(int x) { return parentOrSize[x] < 0 ? x : find(parentOrSize[x]); }
    int time() { return sz(history); }
    void rollback(int checkpoint) {
        for (int i = time(); i-- > checkpoint;) parentOrSize[history[i].first] = history[i].second;
        history.resize(checkpoint);
    }
    bool join(int a, int b) {
        a = find(a), b = find(b);
        if (a == b) return false;
        if (parentOrSize[a] > parentOrSize[b]) swap(a, b);
        history.push_back({a, parentOrSize[a]});
        history.push_back({b, parentOrSize[b]});
        parentOrSize[a] += parentOrSize[b];
        parentOrSize[b] = a;
        return true;
    }
};
```

Matrix

Description: Fixed-size square matrices, initially zero; N is a compile-time dimension. Exponent must be nonnegative and A^0 is the identity. Multiplication has no modulus: add reduction or choose a suitable scalar type, with products fitting that type.
Usage: Matrix<ll,2> a; a.data={{{1,1},{1,0}}};
array<ll,2> state={1,0};
auto next=(a^10)*state; // {89,55}

```
template <class T, int N> struct Matrix {
    typedef Matrix M;
    array<array<T, N>, N> data{};
    M operator*(const M &other) const {
        M result;
        for (int i = 0; i < (N); ++i)
            for (int j = 0; j < (N); ++j)
                for (int k = 0; k < (N); ++k) result.data[i][k] += data[i][j] * other.data[j][k];
        return result;
    }
    array<T, N> operator*(const array<T, N> &vec) const {
        array<T, N> result{};
        for (int i = 0; i < (N); ++i)
            for (int j = 0; j < (N); ++j) result[i] += data[i][j] * vec[j];
        return result;
    }
    M operator^(ll exponent) const {
        assert(exponent >= 0);
        M result, base(*this);
        for (int i = 0; i < (N); ++i) result.data[i][i] = 1;
        while (exponent) {
            if (exponent & 1) result = result * base;
            base = base * base;
            exponent >>= 1;
        }
    }
};
```

```
}
return result;
}
};
```

Convex Hull Trick (CHT)

Description: Add lines of the form slope*x+intercept, and query maximum values at points x. Useful for dynamic programming (“convex hull trick”). Lines and queries may arrive in any order; x, slopes and intercepts are integers. Call add before the first query. For minimum queries, insert (-slope,-intercept) and negate the result. Products, differences and intersection arithmetic must fit ll. lastX marks the final integer x where a line is best.
Usage: CHT hull; hull.add(2,3); hull.add(-1,7);
ll best=hull.query(4); // max(11,3)=11
Time: $\mathcal{O}(\log N)$

```
struct CHTLine {
    mutable ll slope, intercept, lastX;
    bool operator<(const CHTLine &other) const { return slope < other.slope; }
    bool operator<(ll x) const { return lastX < x; }
};
struct CHT : multiset<CHTLine, less<>> {
    // (for doubles, use inf = 1/.0, div(a,b) = a/b)
    static const ll inf = LLONG_MAX;
    ll div(ll a, ll b) { // floored division
        return a / b - ((a ^ b) < 0 && a % b);
    }
    bool isect(iterator line, iterator nextLine) {
        if (nextLine == end()) return line->lastX = inf, 0;
        if (line->slope == nextLine->slope)
            line->lastX = line->intercept > nextLine->intercept ? inf : -inf;
        else
            line->lastX = div(nextLine->intercept - line->intercept, line->slope - nextLine->slope);
        return line->lastX >= nextLine->lastX;
    }
    void add(ll slope, ll intercept) {
        auto nextLine = insert({slope, intercept, 0}), line = nextLine++, prevLine = line;
        while (isect(line, nextLine)) nextLine = erase(nextLine);
        if (prevLine != begin() && isect(--prevLine, line)) isect(prevLine, line = erase(line));
        while ((line = prevLine) != begin() && (--prevLine)->lastX >= line->lastX)
            isect(prevLine, erase(line));
    }
    ll query(ll x) {
        assert(!empty());
        auto line = *lower_bound(x);
        return line.slope * x + line.intercept;
    }
};
```

Implicit Treap

Description: A short self-balancing tree. It acts as a sequential container with log-time splits/joins, and is easy to augment with additional data. This is an implicit treap: in-order position is the key. split(root,k) separates the first k values; merge(a,b) concatenates sequences. Positions and ranges are 0-based, with ranges [l,r). Save roots returned by ins/merge, and recalc after changing children. Extend recalc for sums or lazy tags.
Usage: TreapNode *root=nullptr;
root=ins(root,new TreapNode(7),0);
auto [a,b]=split(root,1); root=merge(a,b);
Time: $\mathcal{O}(\log N)$

```
struct TreapNode {
    TreapNode *left = 0, *right = 0;
```

```
int val, priority, size = 1;
TreapNode(int val) : val(val), priority(rand()) {}
void recalc();
};
int cnt(TreapNode *node) { return node ? node->size : 0; }
void TreapNode::recalc() { size = cnt(left) + cnt(right) + 1; }
template <class F> void each(TreapNode *node, F visit) {
    if (node) {
        each(node->left, visit);
        visit(node->val);
        each(node->right, visit);
    }
}
pair<TreapNode *, TreapNode *> split(TreapNode *node, int
    leftCount) {
    if (!node) return {};
    if (cnt(node->left) >= leftCount) { // "node->val >= key" for
        lower_bound(k)
        auto [leftPart, rightPart] = split(node->left, leftCount);
        node->left = rightPart;
        node->recalc();
        return {leftPart, node};
    } else {
        auto [leftPart, rightPart] =
            split(node->right, leftCount - cnt(node->left) - 1); //
            and just "key"
        node->right = leftPart;
        node->recalc();
        return {node, rightPart};
    }
}
TreapNode *merge(TreapNode *left, TreapNode *right) {
    if (!left) return right;
    if (!right) return left;
    if (left->priority > right->priority) {
        left->right = merge(left->right, right);
        return left->recalc(), left;
    } else {
        right->left = merge(left, right->left);
        return right->recalc(), right;
    }
}
TreapNode *ins(TreapNode *root, TreapNode *node, int pos) {
    auto [left, right] = split(root, pos);
    return merge(merge(left, node), right);
}
// Example application: move the range [left, right) to index
// targetPos
void move(TreapNode *&root, int left, int right, int targetPos)
{
    TreapNode *prefix, *middle, *suffix;
    tie(prefix, middle) = split(root, left);
    tie(middle, suffix) = split(middle, right - left);
    if (targetPos <= left)
        root = merge(ins(prefix, middle, targetPos), suffix);
    else
        root = merge(prefix, ins(suffix, middle, targetPos - right)
        );
}
```

Fenwick Tree (BIT) + Walk

Description: 1-based sums: update(pos,delta) adds rather than assigns; query(pos) sums [1,pos] and query(l,r) is inclusive. lowerBound finds the first prefix reaching a target; it requires nonnegative elements and returns 0 for target<=0 or n+1 if absent. This is Fenwick walking (binary lifting): skip a block only if its whole sum is still below the target. With frequencies, lowerBound(k) locates the kth item, counting duplicates. Updates may be negative, but every resulting element must remain nonnegative for walking to work.

```
Usage: Fenwick bit(5); bit.update(2,3); bit.update(4,7);
11 sum=bit.query(2,4); // 10
int pos=bit.lowerBound(4); // 4
// With 3 items at position 2 and 7 at position 4:
int kth=bit.lowerBound(5); // fifth item is at position 4
// To reach a positive target starting from position 1:
int l=3; 11 target=2;
11 before=bit.query(1-1);
pos=bit.lowerBound(before+target);
// n+1 means insufficient sum in [1,n].
Time:  $\mathcal{O}(\log N)$  per operation;  $\mathcal{O}(N)$  memory.
struct Fenwick {
    int n;
    vector<ll> bit;
    Fenwick(int n = 0) : n(n), bit(n + 1, 0) {}
    void update(int pos, ll delta) {
        for (int i = pos; i <= n; i += i & -i) bit[i] += delta;
    }
    ll query(int pos) const {
        ll sum = 0;
        for (int i = pos; i > 0; i -= i & -i) sum += bit[i];
        return sum;
    }
    ll query(int l, int r) const {
        if (l > r) return 0;
        return query(r) - query(l - 1);
    }
    int lowerBound(ll targetSum) const {
        if (targetSum <= 0) return 0;
        int pos = 0;
        int step = 1;
        while (step <= n / 2) step *= 2;
        // Walk by decreasing powers of two; targetSum is the
        // remaining sum needed.
        for (; step > 0; step /= 2) {
            int next = pos + step;
            if (next <= n && bit[next] < targetSum) {
                // Skip this entire block; no prefix ending inside it
                // can reach the target.
                pos = next;
                targetSum -= bit[next];
            }
        }
        return pos + 1;
    }
};
```

2D Fenwick Tree (BIT)

Description: Offline 2D Fenwick tree. x is in [1, n]; y can be any integer. Register every update coordinate with fakeUpdate before init. query(x,y) sums points with first coordinate <= x and second coordinate <= y. Construct a new object to reset. All future update locations must be registered, even if their eventual values are unknown. Query coordinates need not be registered. Sum a rectangle [x1,x2] by [y1,y2] with four prefix queries using x1-1 and y1-1. Updates add a delta rather than assign a value. Usage: Fenwick2D bit(n); bit.fakeUpdate(2,10); bit.fakeUpdate(4,20); bit.init(); bit.update(2,10,5); 11 sum=bit.query(3,15); // 5 Time: $\mathcal{O}(\log^2 N)$ per operation; $\mathcal{O}(U \log N)$ memory for U registered points.

```
"FenwickTree.h"
struct Fenwick2D {
    int n;
    vector<vector<int>>> yCoords;
    vector<Fenwick> bit;
    Fenwick2D(int n) : n(n), yCoords(n + 1), bit(n + 1) {}
    void fakeUpdate(int x, int y) {
        for (int i = x; i <= n; i += i & -i) yCoords[i].push_back(y
        );
    }
};
```

```

    }
    void init() {
        for (int i = 1; i <= n; i++) {
            sort(all(yCoords[i]));
            yCoords[i].erase(unique(all(yCoords[i])), yCoords[i].end
            ());
            bit[i] = Fenwick(sz(yCoords[i]));
        }
    }
    void update(int x, int y, ll delta) {
        for (int i = x; i <= n; i += i & -i) {
            int pos = lower_bound(all(yCoords[i]), y) - yCoords[i].
            begin() + 1;
            bit[i].update(pos, delta);
        }
    }
    ll query(int x, int y) const {
        ll sum = 0;
        for (int i = x; i > 0; i -= i & -i) {
            int pos = upper_bound(all(yCoords[i]), y) - yCoords[i].
            begin();
            sum += bit[i].query(pos);
        }
        return sum;
    }
};
```

Persistent Fenwick Tree (2D)

Description: Timestamped 2D Fenwick tree: point addition and historical rectangle sums. Coordinates are 1-based; rectangles are inclusive. Initially zero. Updates MUST arrive in nondecreasing time order; same-time additions are combined. Queries may use any time and include all updates at or before it. Only the latest state can be updated; branching is unsupported. Each BIT cell stores time,sum; binary search selects its historical value. The dense N*M history vectors can be expensive. For 1D, remove the y loops and use history[x].

Usage: PersistentFenwick2D bit(n,m); bit.update(2,3,5,1); // add 5 at (2,3), time 1 bit.update(2,3,-2,2); // time 2; negative deltas are allowed 11 old=bit.query(1,1,1,n,m); // rectangle sum at time 1: 5 11 now=bit.query(2,n,m); // prefix rectangle [1,n] x [1,m] at time 2: 3 // Queries before the first update return zero; sums must fit 11. Time: Amortized $\mathcal{O}(\log N \log M)$ update, $\mathcal{O}(\log N \log M \log(U + 1))$ query; $\mathcal{O}(NM + U \log N \log M)$ memory for U updates.

```
struct PersistentFenwick2D {
    int n, m;
    vector<vector<vector<pair<int, ll>>>> history;
    PersistentFenwick2D(int n, int m)
        : n(n), m(m), history(n + 1, vector<vector<pair<int, ll
        >>>(m + 1)) {}
    void update(int x, int y, ll delta, int time) {
        for (int i = x; i <= n; i += i & -i)
            for (int j = y; j <= m; j += j & -j) {
                auto &cell = history[i][j];
                ll value = (cell.empty() ? 0 : cell.back().second) +
                    delta;
                if (!cell.empty() && cell.back().first == time)
                    cell.back().second = value;
                else
                    cell.push_back({time, value});
            }
    }
    ll query(int time, int x, int y) const {
        ll sum = 0;
        for (int i = x; i > 0; i -= i & -i)
            for (int j = y; j > 0; j -= j & -j) {
                const auto &cell = history[i][j];
```

```
        auto it = upper_bound(all(cell), time,
                               [(int t, const pair<int, ll> &
                                entry) { return t < entry.
                                              first; }]);
        if (it != cell.begin()) sum += prev(it)->second;
    }
    return sum;
}
ll query(int time, int x1, int y1, int x2, int y2) const {
    if (x1 > x2 || y1 > y2) return 0;
    return query(time, x2, y1) - query(time, x1 - 1, y2) -
        query(time, x2, y1 - 1) +
        query(time, x1 - 1, y1 - 1);
}
};
```

Sparse Table (RMQ)

Description: Static minimum and maximum on 1-based inclusive [l,r], with l<=r. Input a[0] is ignored. Rebuild after updates; overlapping blocks work for min/max, not sums.

Usage: SparseTable rmq(vector<int>>{0,7,2,5,1});

int minimum=rmq.getMin(2,3); // 2

int maximum=rmq.getMax(2,3); // 5

Time: $\mathcal{O}(N \log N)$ construction and memory; $\mathcal{O}(1)$ per query.

```
struct SparseTable {
    vector<vector<int>> stMin, stMax;
    SparseTable(const vector<int> &a) : stMin(1, a), stMax(1, a) {
        int n = sz(a) - 1;
        for (int k = 1; (1LL << k) <= n; k++) {
            stMin.emplace_back(n + 1);
            stMax.emplace_back(n + 1);
            for (int i = 1; i + (1 << k) - 1 <= n; i++) {
                int j = i + (1 << (k - 1));
                stMin[k][i] = min(stMin[k - 1][i], stMin[k - 1][j]);
                stMax[k][i] = max(stMax[k - 1][i], stMax[k - 1][j]);
            }
        }
    }
    int getMin(int l, int r) const {
        int k = __lg(r - l + 1);
        return min(stMin[k][l], stMin[k][r - (1 << k) + 1]);
    }
    int getMax(int l, int r) const {
        int k = __lg(r - l + 1);
        return max(stMax[k][l], stMax[k][r - (1 << k) + 1]);
    }
};
```

Sliding-Window Minimum and Maximum

Description: Minimum and maximum of every length-k window in a static array. Deques store indices; remove expired indices from the front and dominated values from the back. Equal values keep the newest index. For 1<=k<=n, result[i] describes 0-based [i,i+k-1]. Invalid k returns two empty vectors. Use a monotonic stack instead for nearest smaller/greater queries.

Usage: vector<int> a={4,2,2,5,1};

auto [minimum,maximum]=slidingWindowMinMax(a,3);

// minimum={2,2,1}, maximum={4,5,5}

// For a 1-based array, remove the unused index 0 before calling.

Time: $\mathcal{O}(N)$ time; $\mathcal{O}(k)$ deque memory plus $\mathcal{O}(N)$ output.

```
template <class T> pair<vector<T>, vector<T>>
    slidingWindowMinMax(const vector<T> &a, int k) {
    vector<T> minimum, maximum;
    if (k <= 0 || k > sz(a)) return {minimum, maximum};
    deque<int> minQ, maxQ;
    for (int i = 0; i < sz(a); i++) {
```

RMQ SlidingWindow MoQueries MoTree

```
        while (!minQ.empty() && minQ.front() <= i - k) minQ.
            pop_front();
        while (!maxQ.empty() && maxQ.front() <= i - k) maxQ.
            pop_front();
        while (!minQ.empty() && a[minQ.back()] >= a[i]) minQ.
            pop_back();
        while (!maxQ.empty() && a[maxQ.back()] <= a[i]) maxQ.
            pop_back();
        minQ.push_back(i);
        maxQ.push_back(i);
        if (i >= k - 1) {
            minimum.push_back(a[minQ.front()]);
            maximum.push_back(a[maxQ.front()]);
        }
    }
    return {minimum, maximum};
}
```

Mo’s Algorithm on Arrays

Description: Offline inclusive queries [l,r] on a[1..N]; require 1<=l<=r<=N. Fill add/del/calc and begin with empty state; side=0/1 means left/right endpoint. No value updates between queries; answers follow input order. The included Hilbert helper precomputes sorting keys without block-size tuning. bits covers every endpoint; int indices fit ull keys. A single-element query has l=r.

Usage: auto answers=mo({{1,3},{2,5},{3,3}});

// Fill add/del/calc; [1,3] means positions 1,2,3; [3,3] means only position 3.

Time: $\mathcal{O}(Q \log Q + Q \log(N + 1) + N \sqrt{Q})$ with $\mathcal{O}(1)$ add/del/calc; $\mathcal{O}(Q + \log(N + 1))$ memory excluding custom state.

```
inline ull hilbertOrder(int x, int y, int bits, int rotation = 0) {
    if (bits == 0) return 0;
    int half = 1 << (bits - 1);
    int quadrant = x < half ? (y < half ? 0 : 3) : (y < half ? 1 : 2);
    quadrant = (quadrant + rotation) & 3;
    const int turn[4] = {3, 0, 0, 1};
    int nextRotation = (rotation + turn[quadrant]) & 3;
    ull squareSize = 1ULL << (2 * bits - 2);
    ull offset = hilbertOrder(x & (half - 1), y & (half - 1),
        bits - 1, nextRotation);
    if (quadrant == 0 || quadrant == 3) offset = squareSize - offset - 1;
    return quadrant * squareSize + offset;
}
void add(int pos, int side) { ... } // add a[pos]
void del(int pos, int side) { ... } // remove a[pos]
int calc() {...} // current answer
vector<int> mo(vector<pii> queries) {
    int left = 1, right = 0; // Initially empty; maintained interval is [left,right].
    int maxEndpoint = 0, bits = 0;
    for (auto [l, r] : queries) maxEndpoint = max(maxEndpoint, r)
        ;
    while ((1LL << bits) <= maxEndpoint) bits++;
    vector<int> order(sz(queries)), answers(sz(queries));
    vector<ull> key(sz(queries));
    for (int i = 0; i < sz(queries); i++)
        key[i] = hilbertOrder(queries[i].first, queries[i].second,
            bits);
    iota(all(order), 0);
    sort(all(order), [&](int queryA, int queryB) { return key[
        queryA] < key[queryB]; });
    for (int queryId : order) {
        auto [queryLeft, queryRight] = queries[queryId];
        while (left > queryLeft) add(--left, 0);
        while (right < queryRight) add(++right, 1);
        while (left < queryLeft) del(left++, 0);
```

```
        while (right > queryRight) del(right--, 1);
        answers[queryId] = calc();
    }
    return answers;
}
```

Mo’s Algorithm on Trees

Description: KACTL’s direct endpoint walking with alternating DFS order and snake sorting. Inclusive vertex paths on 1..N; adj has size N+1, both edge directions, and default root=1. Fill add/del/calc; side=0/1 selects the endpoint. Begin with empty state; reset before reuse. No value updates between queries. rank becomes a climb stack after sorting; no LCA table is needed. Both endpoints are included, even for (u,u). Answers follow input order. DFS stack can be linear. For edge values, process the traversed edge in step and omit add(root,0), keeping active[root]. Block size is near N/sqrt(Q); adjust if needed.

Usage: // Fill add/del/calc for the desired statistic (e.g. distinct vertex values).

vector<pii> queries={{u,v},{u,u}};

auto answers=moTree(queries,adj); // 1-based vertices, both

path endpoints included

// For another root in 1..N: moTree(queries,adj,root).

// Compress values first if callbacks index a frequency array by value.

Time: $\mathcal{O}(N + Q \log Q + N \sqrt{Q})$ with $\mathcal{O}(1)$ add/del/calc; $\mathcal{O}(N + Q)$ memory.

```
void add(int u, int side) { ... } // Add vertex u at this end of the path.
```

```
void del(int u, int side) { ... } // Remove vertex u at this end of the path.
```

```
int calc() {...} // Current path answer.
vector<int> moTree(const vector<pii> &queries, const vector<
    vector<int>> &adj, int root = 1) {
    if (queries.empty()) return {};
    int n = sz(adj) - 1, timer = 1, pos[2] = {root, root};
    int blockSize = max(1, (int)(n / sqrt(sz(queries))));
    vector<int> order(sz(queries)), answers(sz(queries));
    vector<int> rank(n + 1), tin(n + 1), tout(n + 1), parent(n + 1);
    vector<bool> active(n + 1);
    add(root, 0);
    active[root] = true;
    auto dfs = [&](int u, int p, bool oddDepth, auto &self) ->
        void {
        parent[u] = p;
        tin[u] = timer;
        if (oddDepth) rank[u] = timer++;
        for (int v : adj[u])
            if (v != p) self(v, u, !oddDepth, self);
        if (!oddDepth) rank[u] = timer++;
        tout[u] = timer - 1; // Inclusive subtree bounds in the alternating DFS order.
    };
    dfs(root, 0, false, dfs);
    iota(all(order), 0);
    sort(all(order), [&](int a, int b) {
        int blockA = rank[queries[a].first] / blockSize;
        int blockB = rank[queries[b].first] / blockSize;
        if (blockA != blockB) return blockA < blockB;
        int rightA = rank[queries[a].second], rightB = rank[queries
            [b].second];
        return blockA & 1 ? rightA > rightB : rightA < rightB;
    });
    for (int id : order) {
        for (int side = 0; side < 2; side++) {
            int &current = pos[side];
            int target = side == 0 ? queries[id].first : queries[id].
                second;
            int climbed = 0;
```

```
auto step = [&](int next) {
    if (active[next]) {
        del(current, side);
        active[current] = false;
    } else {
        add(next, side);
        active[next] = true;
    }
    current = next;
};
// Climb target until it is an ancestor of current; save
// its descent path.
while (!(tin[target] <= tin[current] && tout[current] <=
    tout[target])) {
    rank[++climbed] = target; // DFS ranks are no longer
    // needed after sorting.
    target = parent[target];
}
while (current != target) step(parent[current]);
while (climbed) step(rank[climbed--]);
}
answers[id] = calc();
}
return answers;
}
```

Graph (7)

7.1 Fundamentals

Dial’s Algorithm (Small Integer Weights)

Description: Single-source shortest paths for integer weights in [0,K], with small K. Use K+1 cyclic buckets; current is the real distance, only bucket indices use modulo. Skip stale entries. K=0 and zero-weight edges work; K=1 covers 0-1 BFS. Allocate adj for your indexing; unreachable distances are extttLLONG.MAX. Negative weights are invalid. Large K needs too much scanning/memory; use heap Dijkstra.

Usage: vector<vector<pii>> adj(n+1); // {neighbor,weight}, weights 0..9

adj[u].push_back({v,3}); // add reverse edge too if undirected

auto dist=dial(adj,1,9); // source=1, K=9

if (dist[target]!=LLONG.MAX) cout << dist[target];

// Use dial(adj,source,1) for weights only 0 and 1.

Time: $\mathcal{O}(N + M + NK)$ time; $\mathcal{O}(N + M + K)$ auxiliary memory including queued entries.

```
vector<ll> dial(const vector<vector<pii>> &adj, int source, int
K) {
    int bucketCount = K + 1;
    vector<queue<int>> buckets(bucketCount);
    vector<ll> dist(sz(adj), LLONG_MAX);
    ll current = 0, pending = 1;
    dist[source] = 0;
    buckets[0].push(source);
    while (pending) {
        while (buckets[current % bucketCount].empty()) current++;
        auto &bucket = buckets[current % bucketCount];
        int u = bucket.front();
        bucket.pop();
        pending--;
        if (dist[u] != current) continue; // Ignore an outdated
        // queued distance.
        for (auto [v, weight] : adj[u]) {
            ll nextDist = current + weight;
            if (nextDist >= dist[v]) continue;
            dist[v] = nextDist;
            buckets[nextDist % bucketCount].push(v);
            pending++;
        }
    }
}
```

```
}
return dist;
}
```

Bellman-Ford Shortest Paths

Description: Calculates shortest paths from s in a graph that might have negative edge weights. Unreachable nodes get dist = inf; nodes reachable through negative-weight cycles get dist = -inf. Assumes $V^2 \max |w_i| < \sim 2^{63}$. Use fresh default-initialized nodes for each run. Vertices are 0-based and edges are directed; add both directions for an undirected graph. prev gives a predecessor only for finite shortest paths. The algorithm sorts edges in place; do not reconstruct paths through negative cycles.

Usage: vector<BellmanFordState> nodes(3);

vector<BellmanFordEdge> edges={{0,1,-2},{1,2,5}};

bellmanFord(nodes,edges,0); // nodes[2].dist==3

Time: $\mathcal{O}(VE)$

```
const ll inf = LLONG_MAX;
struct BellmanFordEdge {
    int from, to, weight, sortKey() { return from < to ? from : -
        from; }
};
struct BellmanFordState {
    ll dist = inf;
    int prev = -1;
};
void bellmanFord(vector<BellmanFordState> &nodes, vector<
    BellmanFordEdge> &edges, int source) {
    nodes[source].dist = 0;
    sort(all(edges), [](BellmanFordEdge first, BellmanFordEdge
        second) {
        return first.sortKey() < second.sortKey();
    });

    int rounds = sz(nodes) / 2 + 2; // /3+100 with shuffled
    // vertices
    for (int i = 0; i < (rounds); ++i)
        for (BellmanFordEdge edge : edges) {
            BellmanFordState fromState = nodes[edge.from], &toState =
                nodes[edge.to];
            if (abs(fromState.dist) == inf) continue;
            ll newDist = fromState.dist + edge.weight;
            if (newDist < toState.dist) {
                toState.prev = edge.from;
                toState.dist = (i < rounds - 1 ? newDist : -inf);
            }
        }
    for (int i = 0; i < (rounds); ++i)
        for (BellmanFordEdge edge : edges)
            if (nodes[edge.from].dist == -inf) nodes[edge.to].dist =
                -inf;
}
```

Topological Sort

Description: DAG skeleton with vertices 1..n: define maxN and read n,m first. The stack pops vertices in topological order; ans[u] is the 1-based position of u. A detected cycle prints an error and exits. Clear g and visited between cases.

Usage: // Define maxN and read n,m before the shown main body.

// Reset visited and clear g for another test case.

```
int visited[maxN], ans[maxN];
vector<int> g[maxN];
stack<int> topo; // reversed order of topo;
void dfs(int u) {
    visited[u] = 1;
    for (auto v : g[u]) {
        if (visited[v] == 1) {
            cout << "Error: graph contains a cycle";
            exit(0);
        }
    }
}
```

```
}
if (!visited[v]) dfs(v);
}
topo.push(u);
visited[u] = 2;
}
int main() {
    while (m--) {
        int u, v; cin >> u >> v;
        g[u].push_back(v);
    }
    for (int i = 1; i <= n; ++i)
        if (!visited[i]) dfs(i);
}
```

7.2 Matching

Hopcroft-Karp Matching

Description: Maximum bipartite matching. adj lists right neighbors for each left vertex; the two sides use separate 0-based indices. Initialize matchRight with rightSize entries of -1. Returns matching size; matchRight[v] is its matched left vertex, or -1.

Usage: vector<vector<int>> adj={{0,1},{1}};

vector<int> matchRight(2,-1);

int size=hopcroftKarp(adj,matchRight); // 2

Time: $\mathcal{O}(E\sqrt{V})$

```
int hopcroftKarp(vector<vector<int>> &adj, vector<int> &
matchRight) {
    int n = sz(adj), matchingSize = 0;
    vector<int> matchLeft(n, -1), queue(n), dist(n);
    auto dfs = [&](auto self, int u) -> bool {
        int nextDist = exchange(dist[u], 0) + 1;
        for (int v : adj[u])
            if (matchRight[v] == -1 || (dist[matchRight[v]] ==
                nextDist && self(self, matchRight[v])))
                return matchLeft[u] = v, matchRight[v] = u, 1;
        return 0;
    };
    for (int queueSize = 0, foundPath = 0;; queueSize = foundPath
        = 0, dist.assign(n, 0)) {
        for (int i = 0; i < (n); ++i)
            if (matchLeft[i] == -1) queue[queueSize++] = i, dist[i] =
                1;
        for (int i = 0; i < (queueSize); ++i)
            for (int v : adj[queue[i]])
                if (matchRight[v] == -1)
                    foundPath = 1;
            else if (!dist[matchRight[v]])
                dist[matchRight[v]] = dist[queue[i]] + 1, queue[
                    queueSize++] = matchRight[v];
        if (!foundPath) return matchingSize;
        for (int i = 0; i < (n); ++i)
            if (matchLeft[i] == -1) matchingSize += dfs(dfs, i);
    }
}
```

Bipartite Matching (DFS)

Description: Maximum bipartite matching for smaller graphs. adj lists right neighbors of left vertices, with separate 0-based indices per side. Initialize matchRight to -1; it stores the matched left vertex or -1. Returns matching size; prefer HopcroftKarp when V*E is large.

Usage: vector<vector<int>> adj={{0,1},{1}};

vector<int> matchRight(2,-1);

int size=dfsMatching(adj,matchRight); // 2

Time: $\mathcal{O}(VE)$

```
bool find(int v, vector<vector<int>> &adj, vector<int> &
matchRight, vector<int> &visited) {
    if (matchRight[v] == -1) return 1;
```

```
visited[v] = 1;
int matchedLeft = matchRight[v];
for (int nextRight : adj[matchedLeft])
    if (!visited[nextRight] && find(nextRight, adj, matchRight, visited)) {
        matchRight[nextRight] = matchedLeft;
        return 1;
    }
return 0;
}

int dfsMatching(vector<vector<int>> &adj, vector<int> &matchRight) {
    vector<int> visited;
    for (int i = 0; i < (sz(adj)); ++i) {
        visited.assign(sz(matchRight), 0);
        for (int v : adj[i])
            if (find(v, adj, matchRight, visited)) {
                matchRight[v] = i;
                break;
            }
    }
    return sz(matchRight) - (int)count(all(matchRight), -1);
}
```

Bipartite Vertex Cover
Description: Bipartite only: adj has n left vertices listing right indices 0..m-1. Returned IDs <n are left vertices; IDs n+v denote right vertex v. Vertices outside the cover form a maximum independent set.
Usage: auto vertices=cover(adj,n,m);
// id<n: left id; otherwise right id-n.
"DFSMatching.h"

```
vector<int> cover(vector<vector<int>> &adj, int n, int m) {
    vector<int> matchRight(m, -1);
    int matchingSize = dfsMatching(adj, matchRight);
    vector<bool> reachableLeft(n, true), reachableRight(m);
    for (int matchedLeft : matchRight)
        if (matchedLeft != -1) reachableLeft[matchedLeft] = false;
    vector<int> stack, cover;
    for (int i = 0; i < (n); ++i)
        if (reachableLeft[i]) stack.push_back(i);
    while (!stack.empty()) {
        int i = stack.back();
        stack.pop_back();
        reachableLeft[i] = 1;
        for (int v : adj[i])
            if (!reachableRight[v] && matchRight[v] != -1) {
                reachableRight[v] = true;
                stack.push_back(matchRight[v]);
            }
    }
    for (int i = 0; i < (n); ++i)
        if (!reachableLeft[i]) cover.push_back(i);
    for (int i = 0; i < (m); ++i)
        if (reachableRight[i]) cover.push_back(n + i);
    assert(sz(cover) == matchingSize);
    return cover;
}
```

Hungarian Algorithm (Assignment)
Description: Minimum-cost assignment of every left vertex to a distinct right vertex. cost is a rectangular 0-based N*M matrix with N<=M; returns totalCost,match, where match[i] is the assigned column. Negate costs for maximum cost. Missing edges need a large cost and a check afterward. Use ll if totals may exceed int.
Usage: vector<vector<ll>> costs={{3,1,8},{2,4,6}};
auto [minCost,match]=weightedMatching(costs); // cost=3
// Row i is assigned to column match[i].
Time: $\mathcal{O}(N^2M)$

```
template <class T> pair<T, vector<int>> weightedMatching(vector<vector<T>> &costs) {
    int i = sz(costs), m = i ? sz(costs[0]) : 0, col, scannedCols, row;
    vector<T> dist(m), potential(m);
    vector<int> matchLeft(i), matchRight(m, -1), colOrder(m), parentRow(m);
    T minDist = 0, newDist, cost = 0;
    while (i--) {
        for (int col = 0; col < (m); ++col)
            dist[col] = costs[i][col], colOrder[col] = col, parentRow[col] = i;
        for (scannedCols = 0;;) {
            for (int j = scannedCols; j < (m); ++j) {
                col = colOrder[j], newDist = dist[col] - potential[col];
                if (j == scannedCols || minDist > newDist)
                    minDist = newDist, swap(colOrder[scannedCols], colOrder[j]);
            }
            if ((row = matchRight[col = colOrder[scannedCols++]]) == -1) break;
            for (int j = 0; j < (m); ++j)
                if (dist[j] > (newDist = costs[row][j] - costs[row][col] + dist[col]))
                    dist[j] = newDist, parentRow[j] = row;
        }
        cost += dist[col];
        while (scannedCols--) potential[colOrder[scannedCols]] = dist[colOrder[scannedCols]] - minDist;
        for (; row != i; swap(col, matchLeft[row])) row = matchRight[col] = parentRow[col];
    }
    return {cost, matchLeft};
}
```

General Graph Matching
Description: Maximum-cardinality matching in an undirected general graph (not weighted). Vertices are 0..n-1; list each edge once, with no self-loops. Returns matched pairs. Randomized Tutte-matrix inversion can fail with probability N/mod .
Usage: vector<pii> edges={{0,1},{1,2},{2,0}};
auto matching=generalMatching(3,edges); // one pair
Time: $\mathcal{O}(N^3)$
"../numerical/MatrixInverse-mod.h"

```
vector<pii> generalMatching(int n, vector<pii> &edges) {
    vector<vector<ll>> tutte(n, vector<ll>(n)), inverse;
    for (pii edge : edges)
        int u = edge.first, v = edge.second, randomValue = rand() % mod;
        tutte[u][v] = randomValue, tutte[v][u] = (mod - randomValue) % mod;
    }

    int rank = matInv(inverse = tutte), extendedSize = 2 * n - rank, matchedU, matchedV;
    assert(rank % 2 == 0);

    if (extendedSize != n) do {
        tutte.resize(extendedSize, vector<ll>(extendedSize));
        for (int i = 0; i < (n); ++i) {
            tutte[i].resize(extendedSize);
            for (int j = n; j < (extendedSize); ++j) {
                int randomValue = rand() % mod;
                tutte[i][j] = randomValue, tutte[j][i] = (mod - randomValue) % mod;
            }
        }
    } while (matInv(inverse = tutte) != extendedSize);
}
```

```
vector<int> active(extendedSize, 1);
vector<pii> matching;
for (int it = 0; it < (extendedSize / 2); ++it) {
    for (int i = 0; i < (extendedSize); ++i)
        if (active[i])
            for (int j = i + 1; j < (extendedSize); ++j)
                if (inverse[i][j] && tutte[i][j]) {
                    matchedU = i;
                    matchedV = j;
                    goto done;
                }
    assert(0);
done:
    if (matchedV < n) matching.emplace_back(matchedU, matchedV);
    active[matchedU] = active[matchedV] = 0;
    for (int pass = 0; pass < (2); ++pass) {
        ll pivotInverse = modpow(inverse[matchedU][matchedV], mod - 2);
        for (int i = 0; i < (extendedSize); ++i)
            if (active[i] && inverse[i][matchedV]) {
                ll factor = inverse[i][matchedV] * pivotInverse % mod;
                for (int j = 0; j < (extendedSize); ++j)
                    inverse[i][j] = (inverse[i][j] - inverse[matchedU][j] * factor) % mod;
            }
        swap(matchedU, matchedV);
    }
    return matching;
}
```

7.3 DFS algorithms
Strongly Connected Components (SCC)
Description: Directed SCCs; default vertices 1..n, or pass firstVertex=0. comp[u] is a 0-based component ID and components[id] lists its vertices. Inter-component edges go to smaller IDs; process IDs in descending order for topological order. Build a new object after changing edges.
Usage: SCC scc(adj); bool same=scc.comp[u]==scc.comp[v];
for (int id=sz(scc.components)-1;id>=0;id--) {
// Process scc.components[id] in topological order.
}
Time: $\mathcal{O}(N + E)$ time and memory.

```
struct SCC {
    int timer = 0;
    vector<int> num, low, comp, activeStack;
    vector<bool> inStack;
    vector<vector<int>> components;
    SCC(const vector<vector<int>> &adj, int firstVertex = 1) {
        int n = sz(adj);
        num.assign(n, 0);
        low.assign(n, 0);
        comp.assign(n, -1);
        inStack.assign(n, false);
        for (int u = firstVertex; u < n; u++)
            if (num[u] == 0) dfs(adj, u);
    }
    void dfs(const vector<vector<int>> &adj, int u) {
        timer++;
        num[u] = timer;
        low[u] = timer;
        activeStack.push_back(u);
        inStack[u] = true;
        for (int v : adj[u]) {
            if (num[v] == 0) {
                dfs(adj, v);
            }
        }
    }
}
```

```
        low[u] = min(low[u], low[v]);
    } else if (inStack[v]) {
        low[u] = min(low[u], num[v]);
    }
}
if (low[u] != num[u]) return;
int id = sz(components);
components.push_back({});
while (true) {
    int v = activeStack.back();
    activeStack.pop_back();
    inStack[v] = false;
    comp[v] = id;
    components[id].push_back(v);
    if (v == u) break;
}
}
};
```

Biconnected Components

Description: Finds all biconnected components in an undirected graph, and runs a callback for the edges in each. In a biconnected component there are at least two internally disjoint paths between any two nodes (a cycle exists through them). Note that a node can be in several components. An edge which is not in a component is a bridge, i.e., not part of any cycle. Assign each undirected edge a unique ID and store that same ID at both endpoints, including parallel edges. The callback receives edge IDs, not vertex IDs. This implementation omits bridges; process the marked bridge branch separately if needed. Clear edgeStack and reset timer before a new graph.

Usage: vector<pii> edges={0,1},{1,2},{2,0};
auto colors=edgeColoring(3,edges);
Time: $\mathcal{O}(E+V)$

```
vector<int> num, edgeStack;
vector<vector<pii>> adj;
int timer;
template <class F> int dfs(int u, int parentEdge, F &
    visitComponent) {
    int entryTime = num[u] = ++timer, low = entryTime;
    for (auto [v, edgeId] : adj[u])
        if (edgeId != parentEdge) {
            if (num[v]) {
                low = min(low, num[v]);
                if (num[v] < entryTime) edgeStack.push_back(edgeId);
            } else {
                int stackStart = sz(edgeStack);
                int childLow = dfs(v, edgeId, visitComponent);
                low = min(low, childLow);
                if (childLow == entryTime) {
                    edgeStack.push_back(edgeId);
                    visitComponent(vector<int>(edgeStack.begin() +
                        stackStart, edgeStack.end()));
                    edgeStack.resize(stackStart);
                } else if (childLow < entryTime)
                    edgeStack.push_back(edgeId);
                else { /* e is a bridge */
                }
            }
        }
    return low;
}
template <class F> void bicomps(F visitComponent) {
    num.assign(sz(adj), 0);
    for (int i = 0; i < (sz(adj)); ++i)
        if (!num[i]) dfs(i, -1, visitComponent);
}
```

2-SAT

Description: 0-based variables; literals x and ~x mean true and false. either(a,b) adds a OR b; implication a=>b is either(~a,b). setValue fixes a literal true. Read values only if solve() succeeds. atMostOne can add auxiliary variables.

Usage: TwoSat ts(3); ts.either(0,1); ts.setValue(2);
bool possible=ts.solve();
if (possible) { int valueOfZero=ts.values[0]; }
Time: $\mathcal{O}(N+E)$, where N is the number of boolean variables, and E is the number of clauses.

```
"scc.h"
struct TwoSat {
    int variableCount;
    vector<vector<int>> adj;
    vector<int> values; // 0 = false, 1 = true
    TwoSat(int variableCount = 0) : variableCount(variableCount),
        adj(2 * variableCount) {}
    int addVar() { // (optional)
        adj.emplace_back();
        adj.emplace_back();
        return variableCount++;
    }
    void either(int literalA, int literalB) {
        literalA = max(2 * literalA, -1 - 2 * literalA);
        literalB = max(2 * literalB, -1 - 2 * literalB);
        adj[literalA].push_back(literalB ^ 1);
        adj[literalB].push_back(literalA ^ 1);
    }
    void setValue(int x) { either(x, x); }
    void atMostOne(const vector<int> &literals) { // (optional)
        if (sz(literals) <= 1) return;
        int currentLiteral = ~literals[0];
        for (int i = 2; i < (sz(literals)); ++i) {
            int auxiliaryVar = addVar();
            either(currentLiteral, ~literals[i]);
            either(currentLiteral, auxiliaryVar);
            either(~literals[i], auxiliaryVar);
            currentLiteral = ~auxiliaryVar;
        }
        either(currentLiteral, ~literals[1]);
    }
    bool solve() {
        SCC scc(adj, 0);
        values.assign(variableCount, 0);
        for (int i = 0; i < variableCount; i++) {
            if (scc.comp[2 * i] == scc.comp[2 * i + 1]) return false;
            // Node 2*i means false; node 2*i+1 means true.
            values[i] = scc.comp[2 * i] > scc.comp[2 * i + 1];
        }
        return true;
    }
};
```

Euler Walk

Description: Hierholzer traversal uses every edge exactly once. Vertices are 0-based and edge IDs are 0..edgeCount-1. For undirected edges, add both directions with the same ID; for directed edges, add one entry. For an open trail, choose an odd-degree start in an undirected graph, or a vertex with outdegree=indegree+1 in a directed graph. For a cycle choose a vertex incident to an edge. Returns edgeCount+1 vertices or an empty vector on failure. Open trails have different endpoints.

Usage: vector<vector<pii>> adj(3);
adj[0].push_back({1,0}); adj[1].push_back({2,1});
auto walk=eulerWalk(adj,2,0); // directed: {0,1,2}
Time: $\mathcal{O}(V+E)$

```
vector<int> eulerWalk(vector<vector<pii>> &adj, int edgeCount,
    int source = 0) {
    int n = sz(adj);
```

```
vector<int> balance(n), nextEdge(n), usedEdge(edgeCount),
    walk, stack = {source};
balance[source]++; // to allow Euler paths, not just cycles
while (!stack.empty()) {
    int u = stack.back(), v, edgeId, &it = nextEdge[u], degree
        = sz(adj[u]);
    if (it == degree) {
        walk.push_back(u);
        stack.pop_back();
        continue;
    }
    tie(v, edgeId) = adj[u][it++];
    if (!usedEdge[edgeId]) {
        balance[u]--, balance[v]++;
        usedEdge[edgeId] = 1;
        stack.push_back(v);
    }
}
for (int degreeDelta : balance)
    if (degreeDelta < 0 || sz(walk) != edgeCount + 1) return {}
return {walk.rbegin(), walk.rend()};
}
```

7.4 Coloring

Edge Coloring

Description: Given a simple, undirected graph with max degree D , computes a $(D+1)$ -coloring of the edges such that no neighboring edges share a color. (D -coloring is NP-hard, but can be done for bipartite graphs by repeated matchings of max-degree nodes.) Vertices are 0..n-1; input must have no self-loops or parallel edges. colors[i] is the 0-based color of input edge i. The algorithm recolors alternating chains to free a color at both endpoints. It guarantees at most $D+1$ colors, not the minimum number.

Usage: vector<pii> edges={0,1},{1,2},{2,0};
auto colors=edgeColoring(3,edges);
Time: $\mathcal{O}(NM)$

```
vector<int> edgeColoring(int n, vector<pii> edges) {
    // fanColor holds degree counts first, then the colors used
    // by the fan.
    vector<int> fanColor(n + 1), colors(sz(edges)), fan(n),
        freeColor(n), colorPos;
    for (pii edge : edges) ++fanColor[edge.first], ++fanColor[
        edge.second];
    int u, v, colorCount = *max_element(all(fanColor)) + 1;
    vector<vector<int>> adj(n, vector<int>(colorCount, -1));
    for (pii edge : edges) {
        tie(u, v) = edge;
        fan[0] = v;
        colorPos.assign(colorCount, 0);
        int current = u, lastVertex = u, alternateColor, baseColor
            = freeColor[u], fanSize = 0, i = 0;
        while (alternateColor = freeColor[v],
            !colorPos[alternateColor] && (v = adj[u][
                alternateColor]) != -1)
            colorPos[alternateColor] = ++fanSize, fanColor[fanSize] =
                alternateColor, fan[fanSize] = v;
        fanColor[colorPos[alternateColor]] = baseColor;
        for (int color = alternateColor; current != -1;
            color ^= baseColor ^ alternateColor, current = adj[
                current][color])
            swap(adj[current][color], adj[lastVertex = current][color
                ^ baseColor ^ alternateColor]);
        while (adj[fan[i]][alternateColor] != -1) {
            int left = fan[i], right = fan[++i], rotatedColor =
                fanColor[i];
            adj[u][rotatedColor] = left;
            adj[left][rotatedColor] = u;
            adj[right][rotatedColor] = -1;
```

```
        freeColor[right] = rotatedColor;
    }
    adj[u][alternateColor] = fan[i];
    adj[fan[i]][alternateColor] = u;
    for (int y : {fan[0], u, lastVertex})
        for (int &z = freeColor[y] = 0; adj[y][z] != -1; z++);
}
for (int i = 0; i < (sz(edges)); ++i)
    for (tie(u, v) = edges[i]; adj[u][colors[i]] != v; ++
        colors[i];
    return colors;
}
```

7.5 Heuristics

Maximal Cliques

Description: Runs a callback for all maximal cliques in a graph (given as a symmetric bitset matrix; self-edges not allowed). Callback is given a bitset representing the maximal clique. Maximal means no vertex can be added; it does not mean largest. For a maximum clique use MaximumClique, or examine every callback and keep the largest. Vertices are 0-based and there are at most 128. Pass candidates containing exactly the valid vertex bits; the default mask includes bits beyond n.

Usage: vector adj(n); B candidates; for (int u=0;u<n;u++) candidates.set(u); // Set adj[u][v]=adj[v][u]=1 for each edge. cliques(adj,[] (B clique) { ... },candidates);

Time: $\mathcal{O}\left(3^{n/3}\right)$, much faster for sparse graphs

```
typedef bitset<128> B;
template <class F>
void cliques(vector<B> &adj, F visitClique, B candidates = ~B()
, B excluded = {}, B clique = {}) {
    if (!candidates.any()) {
        if (!excluded.any()) visitClique(clique);
        return;
    }
    auto pivot = (candidates | excluded)._Find_first();
    auto branches = candidates & ~adj[pivot];
    for (int i = 0; i < (sz(adj)); ++i)
        if (branches[i]) {
            clique[i] = 1;
            cliques(adj, visitClique, candidates & adj[i], excluded &
                adj[i], clique);
            clique[i] = candidates[i] = 0;
            excluded[i] = 1;
        }
}
```

Maximum Clique

Description: Quickly finds a maximum clique of a graph (given as symmetric bitset matrix; self-edges not allowed). Can be used to find a maximum independent set by finding a maximum clique of the complement graph. Exact branch-and-bound uses color bounds to prune search while retaining the largest clique. Runtime can still be exponential. Supply a nonempty symmetric adjacency matrix with at most 200 vertices and a zero diagonal. maxClique returns 0-based vertex IDs; construct a fresh object for a new search.

Usage: vb adj(n); // set adj[u][v]=adj[v][u]=1 for each edge MaximumClique solver(adj); auto vertices=solver.maxClique(); **Time:** Runs in about 1s for n=155 and worst case random graphs (p=.90). Runs faster for sparse graphs.

```
typedef vector<bitset<200>> vb;
struct MaximumClique {
    double limit = 0.025, searchCount = 0;
    struct Vertex {
        int id, bound = 0; // degree during init, then a coloring
            upper bound
    };
};
```

```
typedef vector<Vertex> vv;
vb adj;
vv vertices;
vector<vector<int>> colorClasses;
vector<int> bestClique, clique, searchStats, prevStats;
void init(vv &candidates) {
    for (auto &v : candidates) v.bound = 0;
    for (auto &v : candidates)
        for (auto other : candidates) v.bound += adj[v.id][other.
            id];
    sort(all(candidates), [](auto a, auto b) { return a.bound >
        b.bound; });
    int maxDegree = candidates[0].bound;
    for (int i = 0; i < (sz(candidates)); ++i) candidates[i].
        bound = min(i, maxDegree) + 1;
}
void expand(vv &candidates, int level = 1) {
    searchStats[level] += searchStats[level - 1] - prevStats[
        level];
    prevStats[level] = searchStats[level - 1];
    while (sz(candidates)) {
        if (sz(clique) + candidates.back().bound <= sz(bestClique
            )) return;
        clique.push_back(candidates.back().id);
        vv neighbors;
        for (auto v : candidates)
            if (adj[candidates.back().id][v.id]) neighbors.
                push_back({v.id});
        if (sz(neighbors)) {
            if (searchStats[level]++ / ++searchCount < limit) init(
                neighbors);
            int writePos = 0, maxColor = 1, minColor = max(sz(
                bestClique) - sz(clique) + 1, 1);
            colorClasses[1].clear(), colorClasses[2].clear();
            for (auto v : neighbors) {
                int color = 1;
                auto isAdjacent = [&](int i) { return adj[v.id][i]; };
                while (any_of(all(colorClasses[color]), isAdjacent))
                    color++;
                if (color > maxColor) maxColor = color, colorClasses[
                    maxColor + 1].clear();
                if (color < minColor) neighbors[writePos++].id = v.id
                    ;
                colorClasses[color].push_back(v.id);
            }
            if (writePos > 0) neighbors[writePos - 1].bound = 0;
            for (int color = minColor; color < (maxColor + 1); ++
                color)
                for (int i : colorClasses[color])
                    neighbors[writePos].id = i, neighbors[writePos++].
                        bound = color;
            expand(neighbors, level + 1);
        } else if (sz(clique) > sz(bestClique))
            bestClique = clique;
        clique.pop_back(), candidates.pop_back();
    }
}
vector<int> maxClique() {
    init(vertices), expand(vertices);
    return bestClique;
}
MaximumClique(vb adjacency)
: adj(adjacency), colorClasses(sz(adj) + 1), searchStats(
    sz(colorClasses)),
    prevStats(searchStats) {
    for (int i = 0; i < (sz(adj)); ++i) vertices.push_back({i})
        ;
}
```

};

Maximum Independent Set

Description: A maximum independent set in G is a **maximum clique** (largest possible size) in the complement of G . Use MaximumClique. A **maximal clique** only means no vertex can be added; it may be smaller than the maximum, so an arbitrary maximal clique is insufficient. For bipartite graphs, take all vertices outside a minimum vertex cover; see MinimumVertexCover. Complement only pairs of distinct vertices; keep the diagonal zero. The returned clique vertices are the independent-set vertices in the original graph. Maximal independent sets, like maximal cliques, may be much smaller than maximum ones.

Usage: vb complement(n); for (int u=0;u<n;u++) for (int v=0;v<n;v++) complement[u][v]=u!=v && !adj[u][v]; MaximumClique solver(complement); auto independent=solver.maxClique();

7.6 Trees

Binary Lifting + LCA

Description: Binary lifting and LCA in one snippet: up[u][k] is the ancestor 2^k steps above u. Build on a connected tree; root defaults to 1 and is its own parent. goUp needs $0 \leq \text{steps} \leq h[u]$. tin is single-entry DFS order. distance counts edges; weightedDistance sums weights using ll root distances. Rebuild after edge changes; DFS is recursive. Use this for kth ancestors or jump aggregates; Euler-tour + RMQ gives $\mathcal{O}(1)$ LCA queries.

Usage: LCA tree(adj); // vector<vector<int>>, unit edge weights int common=tree.lca(u,v); int ancestor=tree.goUp(u,2); // requires tree.h[u]>=2 int edges=tree.distance(u,v); vector<vector<pll>> weightedAdj(n+1); // Each pair is {neighbor,weight}; add both directions for an undirected tree. LCA weighted(weightedAdj); ll length=weighted.weightedDistance(u,v); **Time:** $\mathcal{O}(N \log N)$ construction and memory; $\mathcal{O}(\log N)$ per query.

```
struct LCA {
    int LOG, timer = 0;
    vector<vector<pll>> adj;
    vector<vector<int>> up;
    vector<int> h, tin;
    vector<ll> rootDist;
    LCA(int n)
        : LOG(__lg(max(1, n)) + 1), adj(n), up(n, vector<int>(LOG
            )), h(n), tin(n), rootDist(n) {}
    LCA(const vector<vector<int>> &graph, int root = 1) : LCA(sz(
        graph)) {
        for (int u = 0; u < sz(graph); u++)
            for (int v : graph[u]) adj[u].push_back({v, 1});
        dfs(root, root);
    }
    LCA(const vector<vector<pll>> &graph, int root = 1) : LCA(sz(
        graph)) {
        adj = graph;
        dfs(root, root);
    }
    void dfs(int u, int parent) {
        tin[u] = timer++;
        up[u][0] = parent;
        for (int k = 1; k < LOG; k++) up[u][k] = up[up[u][k - 1]][k
            - 1];
        for (auto [v, weight] : adj[u]) {
            if (v == parent) continue;
            h[v] = h[u] + 1;
            rootDist[v] = rootDist[u] + weight;
            dfs(v, u);
        }
    }
};
```

```
    }
}
int goUp(int u, int steps) const {
    for (int k = 0; k < LOG; k++)
        if (steps >> k & 1) u = up[u][k];
    return u;
}
int lca(int u, int v) const {
    if (h[u] < h[v]) swap(u, v);
    u = goUp(u, h[u] - h[v]);
    if (u == v) return u;
    for (int k = LOG - 1; k >= 0; k--)
        if (up[u][k] != up[v][k]) u = up[u][k], v = up[v][k];
    return up[u][0];
}
int distance(int u, int v) const { return h[u] + h[v] - 2 * h[lca(u, v)]; }
ll weightedDistance(int u, int v) const {
    return rootDist[u] + rootDist[v] - 2 * rootDist[lca(u, v)];
}
};
```

LCA: Euler Tour + RMQ

Description: Euler-tour + RMQ version of LCA. Record u on entry and again after each child; the LCA is the shallowest vertex between the first occurrences of u and v. This is different from the entry/exit tour used by tree Mo. Build on a connected tree; default root=1. Choose this for many LCA queries. Binary lifting also supports kth ancestors and jump aggregates, with less tour-table overhead. Rebuild after changing edges; DFS is recursive.

Usage: LCAEuler tree(adj); // vector<vector<int>>, undirected or parent-to-child
int common=tree.lca(u,v);
int edges=tree.distance(u,v);
// With separately computed weighted root distances:
// length = rootDist[u] + rootDist[v]
// - 2 * rootDist[common];
Time: $\mathcal{O}(N \log N)$ construction and memory; $\mathcal{O}(1)$ per LCA/distance query.

```
struct LCAEuler {
    vector<int> h, first, tour;
    vector<vector<int>> table;
    LCAEuler(const vector<vector<int>> &adj, int root = 1) : h(sz(adj)), first(sz(adj), -1) {
        dfs(adj, root, root);
        table.push_back(tour);
        int m = sz(tour);
        for (int k = 1; (1LL << k) <= m; k++) {
            table.emplace_back(m - (1 << k) + 1);
            for (int i = 0; i < sz(table[k]); i++)
                table[k][i] = shallower(table[k - 1][i], table[k - 1][i + (1 << (k - 1))]);
        }
    }
    void dfs(const vector<vector<int>> &adj, int u, int parent) {
        first[u] = sz(tour);
        tour.push_back(u);
        for (int v : adj[u]) {
            if (v == parent) continue;
            h[v] = h[u] + 1;
            dfs(adj, v, u);
            tour.push_back(u);
        }
    }
    int shallower(int u, int v) const { return h[u] <= h[v] ? u : v; }
    int lca(int u, int v) const {
        int l = first[u], r = first[v];
        if (l > r) swap(l, r);
        int k = __lg(r - l + 1);
```

```
        return shallower(table[k][l], table[k][r - (1 << k) + 1]);
    }
    int distance(int u, int v) const { return h[u] + h[v] - 2 * h[lca(u, v)]; }
};
```

Virtual Tree Construction

Description: Given a rooted tree and a subset S of nodes, compute the minimal subtree that contains all the nodes by adding all (at most $|S| - 1$) pairwise LCA's and compressing edges. Returns a list of (par, orig_index) representing a tree rooted at virtual index 0. The root points to itself. Use when a query touches only a small subset of a large tree. Build LCA once and compress each subset separately. At virtual index i, result[i].second is the original vertex and result[i].first is its virtual parent index. A compressed edge represents an original path; compute its length from depths or weighted root distances.

Usage: LCA tree(adj);
auto compressed=compressTree(tree,{u,v,w});
for (int i=1;i<sz(compressed);i++) {
 int parentIndex=compressed[i].first;
 int originalVertex=compressed[i].second;
}
Time: $\mathcal{O}(|S|(\log |S| + \log N))$

```
"LCA.h"
typedef vector<pair<int, int>> vpi;
vpi compressTree(LCA &lca, const vector<int> &subset) {
    if (subset.empty()) return {};
    vector<int> virtualIndex;
    virtualIndex.resize(sz(lca.tin));
    vector<int> vertices = subset, &entryTime = lca.tin;
    auto cmp = [&](int a, int b) { return entryTime[a] < entryTime[b]; };
    sort(all(vertices), cmp);
    int m = sz(vertices) - 1;
    for (int i = 0; i < (m); ++i) {
        int a = vertices[i], b = vertices[i + 1];
        vertices.push_back(lca.lca(a, b));
    }
    sort(all(vertices), cmp);
    vertices.erase(unique(all(vertices)), vertices.end());
    for (int i = 0; i < (sz(vertices)); ++i) virtualIndex[vertices[i]] = i;
    vpi virtualTree = {pii(0, vertices[0])};
    for (int i = 0; i < (sz(vertices) - 1); ++i) {
        int a = vertices[i], b = vertices[i + 1];
        virtualTree.emplace_back(virtualIndex[lca.lca(a, b)], b);
    }
    return virtualTree;
}
```

Heavy-Light Decomposition (HLD)

Description: Decomposes a tree into vertex disjoint heavy paths and light edges such that the path from any leaf to the root contains at most $\log(n)$ light edges. Code does additive modifications and max queries, but can support commutative segtree modifications/queries on paths and subtrees. Takes as input the full adjacency list. VALS.EDGES being true means that values are stored in the edges, as opposed to the nodes. All values initialized to the segtree default. Root must be 0. Use for online path updates/queries with a commutative aggregate. The constructor reorders adj in place and assumes vertices 0..n-1 rooted at 0. modifyPath adds a delta; queryPath/-querySubtree return maxima. With edge values each non-root vertex stores its parent edge. Initialize finite tree values before addition, since the default is -inf.

Usage: HLD<false> hld(adj); hld.tree.set(1,sz(adj),0);
hld.modifyPath(u,v,5); int best=hld.queryPath(u,v);
int subtreeBest=hld.querySubtree(u);
Time: $\mathcal{O}((\log N)^2)$

```
"../data-structures/LazySegmentTree.h"
template <bool VALS_EDGES> struct HLD {
```

```
    int n, timer = 0;
    vector<vector<int>> adj;
    vector<int> parent, subtreeSize, head, pos;
    LazySegTree tree;
    HLD(vector<vector<int>> graph) : n(sz(graph)), adj(graph), parent(n, -1), subtreeSize(n, 1), head(n), pos(n), tree(n) {
        dfsSz(0);
        dfsHld(0);
    }
    void dfsSz(int v) {
        for (int &u : adj[v]) {
            adj[u].erase(find(all(adj[u]), v));
            parent[u] = v;
            dfsSz(u);
            subtreeSize[v] += subtreeSize[u];
            if (subtreeSize[u] > subtreeSize[adj[v][0]]) swap(u, adj[v][0]);
        }
    }
    void dfsHld(int v) {
        pos[v] = timer++;
        for (int u : adj[v]) {
            head[u] = (u == adj[v][0] ? head[v] : u);
            dfsHld(u);
        }
    }
    template <class B> void process(int u, int v, B visitRange) {
        for (;;) v = parent[head[v]] {
            if (pos[u] > pos[v]) swap(u, v);
            if (head[u] == head[v]) break;
            visitRange(pos[head[v]], pos[v] + 1);
        }
        visitRange(pos[u] + VALS_EDGES, pos[v] + 1);
    }
    void modifyPath(int u, int v, int val) {
        process(u, v, [&](int l, int r) { tree.add(l + 1, r, val); });
    }
    int queryPath(int u, int v) { // Modify depending on problem
        int answer = INT_MIN;
        process(u, v, [&](int l, int r) { answer = max(answer, tree.query(l + 1, r)); });
        return answer;
    }
    int querySubtree(int v) { // modifySubtree is similar
        return tree.query(pos[v] + VALS_EDGES + 1, pos[v] + subtreeSize[v]);
    }
};
```

Centroid Decomposition

Description: Path-counting skeleton; example counts unordered distinct pairs at exactly k edges. Customize the CHANGE blocks: state/storage, initial/extended state, query/insert/reset. Query a whole child BEFORE inserting it; only the centroid and earlier children are stored. Reset touched entries after each centroid. Recurse to count paths inside one child. For vertex sums, match k+value[centroid]-state; for edge sums adapt adjacency/state extension. Length ranges need a Fenwick range query; XOR needs mask states. Prune only if deeper states cannot become valid again. Root defaults to 1 (pass 0 if needed); use a fresh object per run. Single-vertex paths are excluded. Recursive DFS may have linear stack depth.

Usage: CentroidDecomposition tree(adj,k); // adj: vector<vector<int>>, both directions
ll answer=tree.solve(); // 1-based vertices; exact-k example
// For vertices 0..n-1: tree.solve(0).
// Change State, initialState/extendState, queryState/addState/clearStates for another problem.
// Example: a chain of 4 vertices with k=2 has answer 2.

Time: $\mathcal{O}(N \log N)$ with constant-time state operations; $\mathcal{O}(N)$ auxiliary memory.

```
struct CentroidDecomposition {
    // CHANGE 1: state, parameters, storage.
    using State = int; // Depth; change to ll, mask, or struct.
    int k;
    vector<vector<int>> adj; // Adapt for weighted edges.
    vector<int> subtreeSize, freq, touched;
    vector<bool> removed;
    CentroidDecomposition(const vector<vector<int>> &graph, int
        length)
        : k(length), adj(graph), subtreeSize(sz(graph)), freq(sz(
            graph)), removed(sz(graph)) {}
    // CHANGE 2: initial state and extension.
    State initialState(int centroid) { return 0; }
    State extendState(State state, int v) { return state + 1; }
    // CHANGE 3: match, insert, reset.
    ll queryState(State state, int centroid) {
        ll need = (ll)k - state; // Vertex sums: k + value[centroid]
            - state.
        return 0 <= need && need < sz(freq) ? freq[need] : 0;
    }
    void addState(State state) {
        if (!freq[state]) touched.push_back(state);
        freq[state]++;
    }
    void clearStates() {
        for (int state : touched) freq[state] = 0;
        touched.clear(); // Map: clear; Fenwick: undo updates.
    }
    int countChild(int u, int parent) {
        subtreeSize[u] = 1;
        for (int v : adj[u])
            if (v != parent && !removed[v]) subtreeSize[u] +=
                countChild(v, u);
        return subtreeSize[u];
    }
    int findCentroid(int u, int parent, int total) {
        for (int v : adj[u])
            if (v != parent && !removed[v] && subtreeSize[v] > total
                / 2)
                return findCentroid(v, u, total);
        return u;
    }
    void collectStates(int u, int parent, State state, vector<
        State> &states) {
        states.push_back(state);
        for (int v : adj[u])
            if (v != parent && !removed[v]) collectStates(v, u,
                extendState(state, v), states);
    }
    ll countPaths(int centroid) {
        ll answer = 0;
        State start = initialState(centroid);
        addState(start); // Include the centroid.
        for (int v : adj[centroid]) {
            if (removed[v]) continue;
            vector<State> states;
            collectStates(v, centroid, extendState(start, v), states)
                ;
            for (State state : states) answer += queryState(state,
                centroid);
            for (State state : states) addState(state); // Keep these
                two loops separate!
        }
        clearStates(); // Reset before recursion.
        return answer;
    }
    ll solve(int u = 1) {
```

```
int total = countChild(u, -1);
int centroid = findCentroid(u, -1, total);
ll answer = countPaths(centroid);
removed[centroid] = true;
for (int v : adj[centroid])
    if (!removed[v]) answer += solve(v);
return answer;
}
};
```

7.7 Math

7.7.1 Number of Spanning Trees

Create an $N \times N$ matrix mat , and for each edge $a \rightarrow b \in G$, do $\text{mat}[a][b]--$, $\text{mat}[b][b]++$ (and $\text{mat}[b][a]--$, $\text{mat}[a][a]++$ if G is undirected). Remove the i th row and column and take the determinant; this yields the number of directed spanning trees rooted at i (if G is undirected, remove any row/column).

7.7.2 Erdős–Gallai theorem

A simple graph with node degrees $d_1 \geq \dots \geq d_n$ exists iff $d_1 + \dots + d_n$ is even and for every $k = 1 \dots n$,

$$\sum_{i=1}^k d_i \leq k(k-1) + \sum_{i=k+1}^n \min(d_i, k).$$

Geometry (8)

8.1 Geometric primitives

2D Point and Vector

Description: Use `Point<ll>` for exact integer cross/dot products and `Point<double>` for fractional coordinates. `p.cross(a,b)` means $(a-p) \times (b-p)$. `perp` rotates 90 degrees CCW; `rotate` takes radians. `unit`/`normal` require a nonzero vector; products must fit `T`.
Usage: `Point<ll> a(0,0), b(2,0), c(0,3);`
`ll cross=a.cross(b,c);` // 6, c is left of a->b


```
template <class T> int sgn(T x) { return (x > 0) - (x < 0); }
template <class T> struct Point {
    typedef Point P;
    T x, y;
    explicit Point(T x = 0, T y = 0) : x(x), y(y) {}
    bool operator<(P other) const { return tie(x, y) < tie(other.
        x, other.y); }
    bool operator==(P other) const { return tie(x, y) == tie(
        other.x, other.y); }
    P operator+(P other) const { return P(x + other.x, y + other.
        y); }
    P operator-(P other) const { return P(x - other.x, y - other.
        y); }
    P operator*(T scale) const { return P(x * scale, y * scale);
        }
    P operator/(T scale) const { return P(x / scale, y / scale);
        }
    T dot(P other) const { return x * other.x + y * other.y; }
    T cross(P other) const { return x * other.y - y * other.x; }
    T cross(P a, P b) const { return (a - *this).cross(b - *this)
        ; }
    T dist2() const { return x * x + y * y; }
    double dist() const { return sqrt((double)dist2()); }
    // angle to x-axis in interval [-pi, pi]
    double angle() const { return atan2(y, x); }
```

```
P unit() const { return *this / dist(); } // makes dist()==1
P perp() const { return P(-y, x); } // rotates +90
degrees
P normal() const { return perp().unit(); }
// returns point rotated angle radians ccw around the origin
P rotate(double angle) const {
    return P(x * cos(angle) - y * sin(angle), x * sin(angle) +
        y * cos(angle));
}
friend ostream &operator<<(ostream &os, P other) {
    return os << "(" << other.x << ", " << other.y << ")";
}
};
```

Point-to-Line Distance

Description:

Signed distance to the infinite line $a \rightarrow b$: positive on the left, negative on the right. Require $a \neq b$; take abs for unsigned distance. Cross products must fit the coordinate type.




Usage: `double distance=abs(lineDist(Point<double>(0,0), Point<double>(2,0),Point<double>(1,3)));` // 3
`"Point.h"`

```
template <class P>
double lineDist(P a, P b, P p) {
    return (double)(b - a).cross(p - a) / (b - a).dist();
}
```

Point-to-Segment Distance

Description:

Distance to the closed segment, including endpoints; equal endpoints are supported. Use `Point<double>` and a tolerance when comparing the distance with zero.




Usage: `double distance=segDist(P(0,0),P(2,0),P(3,1));`
`bool on=distance<1e-9;`
`"Point.h"`

```
typedef Point<double> P;
double segDist(P start, P finish, P p) {
    if (start == finish) return (p - start).dist();
    auto length2 = (finish - start).dist2(),
        projection = min(length2, max(.0, (p - start).dot(finish
            - start)));
    return (p - start) * length2 - (finish - start) * projection
        ).dist() / length2;
}
```

Segment Intersection

Description:

If a unique intersection point between the line segments going from s_1 to e_1 and from s_2 to e_2 exists then it is returned. If no intersection point exists an empty vector is returned. If infinitely many exist a vector with 2 elements is returned, containing the endpoints of the common line segment. The wrong position will be returned if P is `Point<ll>` and the intersection point does not have integer coordinates. Products of three coordinates are used in intermediate steps so watch out for overflow if using `int` or long long. Endpoints are included. The returned vector has size 0, 1 or 2 for disjoint segments, one intersection, or an overlapping segment. Use `Point<double>` for fractional coordinates; exact boundary checks may need a tolerance for noisy real input.



Usage: `auto intersection=segInter(Point<double>(0,0), Point<double>(2,2),Point<double>(0,2),Point<double>(2,0));`
// intersection[0] is (1,1).
`"Point.h", "OnSegment.h"`

```
template <class P> vector<P> segInter(P a, P b, P c, P d) {
    auto sideA = c.cross(d, a), sideB = c.cross(d, b), sideC = a.
        cross(b, c), sideD = a.cross(b, d);
```

```
// Checks if intersection is single non-endpoint point.
if (sgn(sideA) * sgn(sideB) < 0 && sgn(sideC) * sgn(sideD) < 0)
    return {(a * sideB - b * sideA) / (sideB - sideA)};
set<P> intersections;
if (onSegment(c, d, a)) intersections.insert(a);
if (onSegment(c, d, b)) intersections.insert(b);
if (onSegment(a, b, c)) intersections.insert(c);
if (onSegment(a, b, d)) intersections.insert(d);
return {all(intersections)};
}
```

Line Intersection

Description: If a unique intersection point of the lines going through s1,e1 and s2,e2 exists {1, point} is returned. If no intersection point exists {0, (0,0)} is returned and if infinitely many exists {-1, (0,0)} is returned. The wrong position will be returned if P is Point<ll> and the intersection point does not have integer coordinates. Products of three coordinates are used in intermediate steps so watch out for overflow if using int or ll. These are infinite lines, not bounded segments; each defining pair must contain distinct points. Inspect the status before reading the point. Use Point<double> for fractional coordinates; integer Point division truncates. **Usage:** auto [status,p]=lineInter(Point<double>(0,0), Point<double>(2,2),Point<double>(0,2),Point<double>(2,0)); // status: 1 unique, 0 parallel, -1 coincident. "Point.h"

```
template <class P> pair<int, P> lineInter(P start1, P end1, P start2, P end2) {
    auto crossDir = (end1 - start1).cross(end2 - start2);
    if (crossDir == 0) // if parallel
        return {-(start1.cross(end1, start2) == 0), P(0, 0)};
    auto weightStart = start2.cross(end1, end2), weightEnd = start2.cross(end2, start1);
    return {1, (start1 * weightStart + end1 * weightEnd) / crossDir};
}
```

Orientation Test (Cross Product)

Description: Returns +1/0/-1 for p left/on/right of start->finish. The eps overload treats distance<=eps as on the line and requires distinct endpoints. Products must fit the coordinate type. **Usage:** int side=sideOf(Point<ll>(0,0),Point<ll>(2,0), Point<ll>(1,3)); // +1: left "Point.h"

```
template <class P> int sideOf(P start, P finish, P p) { return sgn(start.cross(finish, p)); }
template <class P> int sideOf(P start, P finish, P p, double eps) {
    auto crossValue = (finish - start).cross(p - start);
    double tolerance = (finish - start).dist() * eps;
    return (crossValue > tolerance) - (crossValue < -tolerance);
}
```


Point on Segment

Description: Exact membership on a closed segment, including endpoints and zero-length segments. Prefer integer coordinates; for Point<double> compare segDist with a tolerance. **Usage:** bool on=onSegment(Point<ll>(0,0),Point<ll>(2,0), Point<ll>(1,0)); // true "Point.h"

```
template <class P> bool onSegment(P start, P finish, P p) {
    return p.cross(start, finish) == 0 && (start - p).dot(finish - p) <= 0;
}
```

Linear Transformation

Description: Apply the linear transformation (translation, rotation and scaling) which takes line p0-p1 to line q0-q1 to point r.



Use when a directed segment is mapped to another by translation, rotation and uniform scaling. The same mapping is applied to point. The source segment must have nonzero length. This does not implement reflection or arbitrary shear.

Usage: P mapped=linearTransformation(P(0,0),P(1,0), P(2,3),P(2,5),P(1,1)); // (0,5) "Point.h"

```
typedef Point<double> P;
P linearTransformation(P sourceStart, P sourceEnd, P targetStart, P targetEnd, P point) {
    P sourceDir = sourceEnd - sourceStart, targetDir = targetEnd - targetStart,
    transform(sourceDir.cross(targetDir), sourceDir.dot(targetDir));
    return targetStart +
        P((point - sourceStart).cross(transform), (point - sourceStart).dot(transform)) / sourceDir.dist2();
}
```

Angle Ordering

Description: A class for ordering angles (as represented by int points and a number of rotations around the origin). Useful for rotational sweeping. Sometimes also represents points or vectors. Use nonzero integer direction vectors to sort events without atan2. turns distinguishes directions after complete rotations. t90/t180/t360 rotate counterclockwise; segmentAngles returns the shorter angular interval covering two directions. Intermediate integer products must fit their types. **Usage:** vector<Angle> directions={{1,0},{0,1},{-1,0}}; sort(all(directions)); auto quarterTurn=Angle(1,0).t90();

```
struct Angle {
    int x, y;
    int turns;
    Angle(int x, int y, int turns = 0) : x(x), y(y), turns(turns) {}
    Angle operator-(Angle b) const { return {x - b.x, y - b.y, turns}; }
    int half() const { assert(x || y); return y < 0 || (y == 0 && x < 0); }
}
Angle t90() const { return {-y, x, turns + (half() && x >= 0)}; }
Angle t180() const { return {-x, -y, turns + half(); }
Angle t360() const { return {x, y, turns + 1}; }
};
bool operator<(Angle a, Angle b) {
    // add a.dist2() and b.dist2() to also compare distances
    return make_tuple(a.turns, a.half(), a.y * (ll)b.x) <
        make_tuple(b.turns, b.half(), a.x * (ll)b.y);
}
// Given two points, this calculates the smallest angle between them, i.e., the angle that covers the defined line segment.
pair<Angle, Angle> segmentAngles(Angle a, Angle b) {
    if (b < a) swap(a, b);
    return (b < a.t180() ? make_pair(a, b) : make_pair(b, a.t360()));
}
Angle operator+(Angle a, Angle b) { // point a + vector b
    Angle result(a.x + b.x, a.y + b.y, a.turns);
    if (a.t180() < result) result.turns--;
    return result.t180() < a ? result.t360() : result;
}
```

Angle angleDiff(Angle a, Angle b) { // angle b - angle a
 int turnDiff = b.turns - a.turns;
 a.turns = b.turns;
 return {a.x * b.x + a.y * b.y, a.x * b.y - a.y * b.x, turnDiff - (b < a)};
}

8.2 Circles

Circle Intersection

Description: Computes the pair of points at which two circles intersect. Returns false in case of no intersection. On success the pair contains both intersections; at tangency they are equal. Coincident circles have infinitely many intersections and trigger an assertion, so handle that case before calling. Radii must be nonnegative.

Usage: pair<P,P> points; bool intersects=circleInter(P(0,0),P(2,0),2,2, &points); "Point.h"

```
typedef Point<double> P;
bool circleInter(P center1, P center2, double r1, double r2, pair<P, P> *intersections) {
    if (center1 == center2) {
        assert(r1 != r2);
        return false;
    }
    P centerDir = center2 - center1;
    double distance2 = centerDir.dist2(), radiusSum = r1 + r2, radiusDiff = r1 - r2,
    projection = (distance2 + r1 * r1 - r2 * r2) / (distance2 * 2),
    height2 = r1 * r1 - projection * projection * distance2;
    if (radiusSum * radiusSum < distance2 || radiusDiff * radiusDiff > distance2) return false;
    P basePoint = center1 + centerDir * projection,
    offset = centerDir.perp() * sqrt(fmax(0, height2) / distance2);
    *intersections = {basePoint + offset, basePoint - offset};
    return true;
}
```

Circle Tangents

Description: Finds the external tangents of two circles, or internal if r2 is negated. Can return 0, 1, or 2 tangents - 0 if one circle contains the other (or overlaps it, in the internal case, or if the circles are the same); 1 if the circles are tangent to each other (in which case .first = .second and the tangent line is perpendicular to the line between the centers). .first and .second give the tangency points at circle 1 and 2 respectively. To find the tangents of a circle with a point set r2 to 0. Each returned pair is a tangent segment joining the two tangency points. Use Point<double>; negating the second radius requests internal tangents rather than a physically negative radius. Coincident centers return no tangents.

Usage: auto external=tangents(Point<double>(0,0),1, Point<double>(4,0),1); // two tangent pairs "Point.h"

```
template <class P> vector<pair<P, P>> tangents(P c1, double r1, P c2, double r2) {
    P centerDir = c2 - c1;
    double radiusDiff = r1 - r2, distance2 = centerDir.dist2(), height2 = distance2 - radiusDiff * radiusDiff;
    if (distance2 == 0 || height2 < 0) return {};
    vector<pair<P, P>> tangentPairs;
    for (double sign : {-1, 1}) {
        P normal = (centerDir * radiusDiff + centerDir.perp() * sqrt(height2) * sign) / distance2;
        tangentPairs.push_back({c1 + normal * r1, c2 + normal * r2});
    }
    if (height2 == 0) tangentPairs.pop_back();
}
```

```
    return tangentPairs;
}
```

Circle-Polygon Intersection


Description: Returns the area of the intersection of a circle with a ccw polygon. The polygon may be concave, but must be simple with nonzero-length edges. It is copied and left unchanged. The area is signed: counterclockwise input gives positive area; take abs if needed. Radius must be positive.
Usage: double area=abs(circlePoly(P(0,0),1,polygon));
Time: $\mathcal{O}(n)$

```

".../..content/geometry/Point.h"
typedef Point<double> P;
#define arg(p, q) atan2(p.cross(q), p.dot(q))
double circlePoly(P center, double radius, vector<P> polygon) {
    auto triangleArea = [&](P p, P q) {
        auto halfRadius2 = radius * radius / 2;
        P direction = q - p;
        auto linearTerm = direction.dot(p) / direction.dist2(),
            constantTerm = (p.dist2() - radius * radius) /
                direction.dist2();
        auto discriminant = linearTerm * linearTerm - constantTerm;
        if (discriminant <= 0) return arg(p, q) * halfRadius2;
        auto enter = max(0., -linearTerm - sqrt(discriminant)),
            exit = min(1., -linearTerm + sqrt(discriminant));
        if (exit < 0 || 1 <= enter) return arg(p, q) * halfRadius2;
        P intersection1 = p + direction * enter, intersection2 = q
            + direction * (exit - 1);
        return arg(p, intersection1) * halfRadius2 + intersection1.
            cross(intersection2) / 2 +
            arg(intersection2, q) * halfRadius2;
    };
    auto area = 0.0;
    for (int i = 0; i < (sz(polygon)); ++i)
        area += triangleArea(polygon[i] - center, polygon[(i + 1) %
            sz(polygon)] - center);
    return area;
}
```

Circumcircle

Description: The circumcirle of a triangle is the circle intersecting all three vertices. ccRadius returns the radius of the circle going through points A, B and C and ccCenter returns the center of the same circle.



The points must be distinct and noncollinear. This is the circle through the vertices, which is not always the triangle's minimum enclosing circle: an obtuse triangle uses its longest side as a diameter. Use Point<double> for the center.
Usage: P a(0,0), b(2,0), c(0,2);
P center=ccCenter(a,b,c); double radius=ccRadius(a,b,c);
"Point.h"

```

typedef Point<double> P;
double ccRadius(P A, P B, P C) {
    return (B - A).dist() * (C - B).dist() * (A - C).dist() / abs
        ((B - A).cross(C - A)) / 2;
}
P ccCenter(P A, P B, P C) {
    P dirAC = C - A, dirAB = B - A;
    return A + (dirAC * dirAB.dist2() - dirAB * dirAC.dist2()) *
        perp() / dirAC.cross(dirAB) / 2;
}
```

Minimum Enclosing Circle

Description: Computes the minimum circle that encloses a set of points. Input must contain at least one point. Returns (center,radius), not radius squared. Random shuffling makes incremental construction fast on average, and the input is copied. Uses floating tolerance to decide whether a point lies outside.

```

Usage: auto [center,radius]=mec(vector<P>{P(0,0),P(2,0)});
// center=(1,0), radius=1.
Time: expected  $\mathcal{O}(n)$ 
"circumcircle.h"
pair<P, double> mec(vector<P> points) {
    shuffle(all(points), mt19937(time(0)));
    P center = points[0];
    double radius = 0, tolerance = 1 + 1e-8;
    for (int i = 0; i < (sz(points)); ++i)
        if ((center - points[i]).dist() > radius * tolerance) {
            center = points[i], radius = 0;
            for (int j = 0; j < (i); ++j)
                if ((center - points[j]).dist() > radius * tolerance) {
                    center = (points[i] + points[j]) / 2;
                    radius = (center - points[i]).dist();
                    for (int k = 0; k < (j); ++k)
                        if ((center - points[k]).dist() > radius *
                            tolerance) {
                            center = ccCenter(points[i], points[j], points[k]
                                );
                            radius = (center - points[i]).dist();
                        }
                }
            }
    return {center, radius};
}
```

8.3 Polygons

Point in Polygon

Description: Returns true if p lies within the polygon. If strict is true, it returns false for points on the boundary. The algorithm uses products in intermediate steps so watch out for overflow. Works for a simple, possibly concave polygon in either orientation; do not repeat the first vertex at the end. strict=true excludes the boundary. Use integer coordinates for exact boundary decisions, or adapt the on-segment check to a tolerance.
Usage: vector<Point<ll>> polygon={{0,0},{4,0},{0,4}};
bool inside=inPolygon(polygon,Point<ll>(1,1));
Time: $\mathcal{O}(n)$
"Point.h", "OnSegment.h", "SegmentDistance.h"

```

template <class P> bool inPolygon(vector<P> &polygon, P point,
    bool strict = true) {
    int inside = 0, n = sz(polygon);
    for (int i = 0; i < (n); ++i) {
        P nextPoint = polygon[(i + 1) % n];
        if (onSegment(polygon[i], nextPoint, point)) return !strict;
        ;
        //or: if (segDist(p[i], q, a) <= eps) return !strict;
        inside ^=
            ((point.y < polygon[i].y) - (point.y < nextPoint.y)) *
            point.cross(polygon[i], nextPoint) >
                0;
    }
    return inside;
}
```

Polygon Area

Description: Returns twice the signed area: positive CCW, negative clock-wise. Use ordered boundary vertices, at least one point, and no repeated final vertex. Geometric area is abs(area2)/2.0; cross-product sums must fit T.
Usage: vector<Point<ll>> p={{0,0},{4,0},{0,3}};
double area=abs(polygonArea2(p))/2.0; // 6
"Point.h"

```

template <class T> T polygonArea2(vector<Point<T>> &polygon) {
    T area2 = polygon.back().cross(polygon[0]);
    for (int i = 0; i < (sz(polygon) - 1); ++i) area2 += polygon[
        i].cross(polygon[i + 1]);
    return area2;
}
```


```

}

Polygon Center
Description: Area centroid of a simple polygon with ordered boundary ver-
tices and nonzero area; either orientation works. Zero-area input divides by
zero.
Usage: vector<P> p={{0,0},{3,0},{0,3}};
P center=polygonCenter(p); // (1,1)
Time:  $\mathcal{O}(n)$ 
"Point.h"
typedef Point<double> P;
P polygonCenter(const vector<P> &polygon) {
    P weightedCenter(0, 0);
    double area2 = 0;
    for (int i = 0, j = sz(polygon) - 1; i < sz(polygon); j = i
        ++i) {
        weightedCenter = weightedCenter + (polygon[i] + polygon[j])
            * polygon[j].cross(polygon[i]);
        area2 += polygon[j].cross(polygon[i]);
    }
    return weightedCenter / area2 / 3;
}
```

Polygon Cut

Description: Returns a vector with the vertices of a polygon with every-thing to the left of the line going from s to e cut away. Use a convex polygon when a single output polygon is required. The implementation keeps the strictly right side of start->finish, discarding points exactly on the line; it inserts crossing points. The cutting line must have distinct endpoints. The result may be empty.
Usage: auto clipped=polygonCut(polygon,P(0,0),P(1,0));
// Keep y<0 plus crossing points.
"Point.h"



```

typedef Point<double> P;
vector<P> polygonCut(const vector<P> &polygon, P start, P
    finish) {
    vector<P> clipped;
    for (int i = 0; i < (sz(polygon)); ++i) {
        P cur = polygon[i], prev = i ? polygon[i - 1] : polygon.
            back();
        auto sideCurrent = start.cross(finish, cur), sidePrevious =
            start.cross(finish, prev);
        if ((sideCurrent < 0) != (sidePrevious < 0))
            clipped.push_back(cur + (prev - cur) * (sideCurrent / (
                sideCurrent - sidePrevious)));
        if (sideCurrent < 0) clipped.push_back(cur);
    }
    return clipped;
}
```

Convex Hull

Description: Returns extreme vertices in CCW order, omitting interior collinear boundary points. Input is copied and sorted; empty/single-point inputs work, and collinear input reduces to endpoints. Cross products must fit ll.
Usage: vector<P> points={{0,0},{2,0},{0,2},{1,0}};
auto hull=convexHull(points); // (1,0) is omitted
Time: $\mathcal{O}(n \log n)$
"Point.h"



```

typedef Point<ll> P;
vector<P> convexHull(vector<P> points) {
    if (sz(points) <= 1) return points;
    sort(all(points));
    vector<P> hull(sz(points) + 1);
    int chainStart = 0, hullSize = 0;
    for (int pass = 2; pass--> 0; chainStart = --hullSize, reverse(
        all(points)))
```

```
for (P p : points) {
    while (hullSize >= chainStart + 2 && hull[hullSize - 2].
        cross(hull[hullSize - 1], p) <= 0)
        hullSize--;
    hull[hullSize++] = p;
}
return {hull.begin(), hull.begin() + hullSize - (hullSize ==
    2 && hull[0] == hull[1])};
}
```

Convex Hull Diameter

Description: Returns the two points with max distance on a convex hull (ccw, no duplicate/collinear points). Use rotating calipers after convexHull. The hull must be nonempty; a single point returns that point twice. Returns endpoints, so compute dist2 of their difference for an exact squared diameter. Do not pass an unordered point cloud.

```
Usage: auto hull=convexHull(points); auto
ends=hullDiameter(hull);
ll distance2=(ends[0]-ends[1]).dist2();
Time: O(n)
"Point.h"
```

```
typedef Point<ll> P;
array<P, 2> hullDiameter(vector<P> hull) {
    int n = sz(hull), j = n < 2 ? 0 : 1;
    pair<ll, array<P, 2>> best({0, {hull[0], hull[0]}});
    for (int i = 0; i < (j); ++i)
        for (; j = (j + 1) % n; {
            best = max(best, {(hull[i] - hull[j]).dist2(), {hull[i],
                hull[j]}});
            if ((hull[(j + 1) % n] - hull[j]).cross(hull[i + 1] -
                hull[i]) >= 0) break;
        }
    return best.second;
}
```

Point in Convex Polygon

Description: Determine whether a point t lies inside a convex hull (CCW order, with no collinear points). Returns true if point lies within the hull. If strict is true, points on the boundary aren't included. The hull must be nonempty, convex, and in boundary order, without interior collinear vertices; convexHull produces suitable input. Use for many containment queries rather than scanning every edge. strict=false includes edges and vertices.

```
Usage: auto hull=convexHull(points);
bool inside=inHull(hull,Point<ll>(2,3),false);
Time: O(log N)
"Point.h", "sideOf.h", "OnSegment.h"
```

```
typedef Point<ll> P;
bool inHull(const vector<P> &hull, P p, bool strict = true) {
    int lo = 1, hi = sz(hull) - 1, boundaryAllowed = !strict;
    if (sz(hull) < 3) return boundaryAllowed && onSegment(hull
        [0], hull.back(), p);
    if (sideOf(hull[0], hull[lo], hull[hi]) > 0) swap(lo, hi);
    if (sideOf(hull[0], hull[lo], p) >= boundaryAllowed ||
        sideOf(hull[0], hull[hi], p) <= -boundaryAllowed)
        return false;
    while (abs(lo - hi) > 1) {
        int mid = (lo + hi) / 2;
        (sideOf(hull[0], hull[mid], p) > 0 ? hi : lo) = mid;
    }
    return sgn(hull[lo].cross(hull[hi], p)) < boundaryAllowed;
}
```

Line-Convex Hull Intersection

Description: Line-convex polygon intersection. The polygon must be ccw and have no collinear points. lineHull(line, poly) returns a pair describing the intersection of a line with the polygon: • (−1, −1) if no collision, • (i, −1) if touching the corner i, • (i, i) if along side (i, i + 1), • (i, j) if crossing sides (i, i + 1) and (j, j + 1). In the last case, if a corner i is crossed, this is treated as happening on side (i, i + 1). The points are returned in the same order as the line hits the polygon. extrVertex returns the index of the hull point with maximum projection onto a direction. Use a nondegenerate infinite line and a strictly convex counterclockwise hull with at least three vertices. Results are edge/vertex indices, not coordinates; wrap i+1 modulo hull.size(). Intersect reported sides with the line to get coordinates. extrVertex returns the index maximizing dot(direction,point).

```
Usage: auto hit=lineHull(Point<ll>(0,0),Point<ll>(1,0),hull);
int rightmost=extrVertex(hull,Point<ll>(1,0));
Time: O(log n)
"Point.h"
```

```
#define cmp(i, j) sgn(direction.perp().cross(hull[(i)%n]-hull[(j)%n]))
#define extr(i) cmp(i + 1, i) >= 0 && cmp(i, i - 1 + n) < 0
template <class P> int extrVertex(vector<P> &hull, P direction)
{
    int n = sz(hull), lo = 0, hi = n;
    if (extr(0)) return 0;
    while (lo + 1 < hi) {
        int mid = (lo + hi) / 2;
        if (extr(mid)) return mid;
        int loTrend = cmp(lo + 1, lo), midTrend = cmp(mid + 1, mid)
            ;
        (loTrend < midTrend || (loTrend == midTrend && loTrend ==
            cmp(lo, mid))) ? hi : lo) = mid;
    }
    return lo;
}
#define cmpL(i) sgn(a.cross(hull[i], b))
template <class P> array<int, 2> lineHull(P a, P b, vector<P> &
    hull) {
    int maxVertex = extrVertex(hull, (a - b).perp());
    int minVertex = extrVertex(hull, (b - a).perp());
    if (cmpL(maxVertex) < 0 || cmpL(minVertex) > 0) return {-1,
        -1};
    array<int, 2> intersections;
    for (int i = 0; i < (2); ++i) {
        int lo = minVertex, hi = maxVertex, n = sz(hull);
        while ((lo + 1) % n != hi) {
            int mid = ((lo + hi + (lo < hi ? 0 : n)) / 2) % n;
            (cmpL(mid) == cmpL(minVertex) ? lo : hi) = mid;
        }
        intersections[i] = (lo + !cmpL(hi)) % n;
        swap(maxVertex, minVertex);
    }
    if (intersections[0] == intersections[1]) return {
        intersections[0], -1};
    if (!cmpL(intersections[0]) && !cmpL(intersections[1]))
        switch ((intersections[0] - intersections[1] + sz(hull) +
            1) % sz(hull)) {
        case 0:
            return {intersections[0], intersections[0]};
        case 2:
            return {intersections[1], intersections[1]};
        }
    return intersections;
}
```

8.4 Misc. Point Set Problems

Closest Pair

Description: Finds the closest pair of points. Input must contain at least two points. Returns the endpoints rather than their distance; it copies and sorts the input internally. Duplicate points give distance zero. Integer squared distances must fit ll.

```
Usage: auto [a,b]=closest(points);
ll distance2=(a-b).dist2();
Time: O(n log n)
"Point.h"
```

```
typedef Point<ll> P;
pair<P, P> closest(vector<P> points) {
    assert(sz(points) > 1);
    set<P> active;
    sort(all(points), [](P a, P b) { return a.y < b.y; });
    pair<ll, pair<P, P>> best{LLONG_MAX, {P(), P()}};
    int j = 0;
    for (P p : points) {
        P window{1 + (ll)sqrt(best.first), 0};
        while (points[j].y <= p.y - window.x) active.erase(points[j
            ++]);
        auto lo = active.lower_bound(p - window), hi = active.
            upper_bound(p + window);
        for (; lo != hi; ++lo) best = min(best, {(x*lo - p).dist2(),
            {x*lo, p}});
        active.insert(p);
    }
    return best.second;
}
```

KD-Tree Nearest Neighbour

Description: KD-tree (2d, can be extended to 3d) Static nearest-neighbor queries on a nonempty set of integer points. nearest returns (squared distance,point); the query point itself is allowed if present. Bounding boxes prune search, but worst-case queries can visit every point. Uncomment the leaf check to exclude equal points. There are no insertions/deletions.

```
Usage: KDTree tree(points);
auto [distance2,nearestPoint]=tree.nearest(P(3,7));
"Point.h"
```

```
typedef long long T;
typedef Point<T> P;
const T INF = numeric_limits<T>::max();
bool on_x(P a, P b) { return a.x < b.x; }
bool on_y(P a, P b) { return a.y < b.y; }
struct KDNode {
    P point; // if
    this is a leaf, the single point in it
    T minX = INF, maxX = -INF, minY = INF, maxY = -INF; // bounds
    KDNode *first = 0, *second = 0;
    T distance(P p) { // min squared distance to a point
        T closestX = (p.x < minX ? minX : p.x > maxX ? maxX : p.x);
        T closestY = (p.y < minY ? minY : p.y > maxY ? maxY : p.y);
        return (P(closestX, closestY) - p).dist2();
    }
    KDNode(vector<P> &points) : point(points[0]) {
        for (P p : points) {
            minX = min(minX, p.x);
            maxX = max(maxX, p.x);
            minY = min(minY, p.y);
            maxY = max(maxY, p.y);
        }
        if (points.size() > 1) {
            // split on x if width >= height (not ideal...)
            sort(all(points), maxX - minX >= maxY - minY ? on_x :
                on_y);
            // divide by taking half the array for each child (not
            // best performance with many duplicates in the middle)
            int half = sz(points) / 2;
```

```
        first = new KNode({points.begin(), points.begin() + half
        });
        second = new KNode({points.begin() + half, points.end()})
        );
    }
};
struct KDTree {
    KNode *root;
    KDTree(const vector<P> &points) : root(new KNode({all(points
    )))) {}
    pair<T, P> search(KNode *node, P p) {
        if (!node->first) {
            // uncomment if we should not find the point itself:
            // if (p == node->point) return {INF, P()};
            return make_pair((p - node->point).dist2(), node->point);
        }

        KNode *nearChild = node->first, *farChild = node->second;
        T nearDistance = nearChild->distance(p), farDistance =
            farChild->distance(p);
        if (nearDistance > farDistance) swap(farDistance,
            nearDistance), swap(nearChild, farChild);

        // search closest side first, other side if needed
        auto best = search(nearChild, p);
        if (farDistance < best.first) best = min(best, search(
            farChild, p));
        return best;
    }
    // find nearest point to a point, and its squared distance
    // (requires an arbitrary operator< for Point)
    pair<T, P> nearest(P p) { return search(root, p); }
};
```

Strings (9)

Trie (Pointers)

Description: Pointer trie for lowercase a..z, with duplicate strings. end-Count counts strings ending at a node; passCount counts strings with that prefix (including exact matches). addString inserts one copy; deleteString removes one copy, doing nothing if absent. Unused nodes are freed during deletion; the destructor frees the remaining trie. Empty strings are supported; countPrefix("") counts all inserted copies. Change the alphabet size and character mapping for other symbols. Recursive deletion/destruction may use stack depth equal to the longest string. Copying is disabled because the trie owns its pointers.

Usage: Trie trie;
trie.addString("abc");
trie.addString("abc");
trie.addString("abd");
bool found=trie.findString("abc"); // true; a prefix alone is not an exact match
int count=trie.countPrefix("ab"); // 3, including duplicates
trie.deleteString("abc"); // one copy remains
// Use a fresh object per test case; its destructor releases the nodes.

Time: O(L) per operation for string length L; O(26S) memory for S distinct prefixes.

```
struct Trie {
    struct Node {
        Node *child[26] = {};
        int endCount = 0, passCount = 0;
        ~Node() {
            for (Node *next : child) delete next;
        }
    };
    Node *root = new Node();
    Trie() = default;
```

```
Trie(const Trie &) = delete;
Trie &operator=(const Trie &) = delete;
~Trie() { delete root; }
void addString(const string &s) {
    Node *node = root;
    node->passCount++;
    for (char ch : s) {
        int letter = ch - 'a';
        if (!node->child[letter]) node->child[letter] = new Node
        ();
        node = node->child[letter];
        node->passCount++;
    }
    node->endCount++;
}
Node *findNode(const string &s) const {
    Node *node = root;
    for (char ch : s) {
        node = node->child[ch - 'a'];
        if (!node) return nullptr;
    }
    return node;
}
bool findString(const string &s) const {
    Node *node = findNode(s);
    return node && node->endCount;
}
int countPrefix(const string &s) const {
    Node *node = findNode(s);
    return node ? node->passCount : 0;
}
void deleteString(Node *&node, const string &s, int pos) {
    node->passCount--;
    if (pos == sz(s))
        node->endCount--;
    else
        deleteString(node->child[s[pos] - 'a'], s, pos + 1);
    if (node != root && node->passCount == 0) {
        delete node;
        node = nullptr; // Detach the deleted child from its
        parent.
    }
}
void deleteString(const string &s) {
    if (findString(s)) deleteString(root, s, 0);
}
};
```

Knuth-Morris-Pratt (KMP)

Description: kmpSearch(text,pattern) returns all 0-based match starts, including overlaps. Empty pattern matches positions 0..text.size(); no separator is needed. pi[i] from prefixFunction is the longest proper border of text[0..i].

Usage: auto positions=kmpSearch("ababa","aba"); // {0,2}
auto pi=prefixFunction("abab"); // {0,0,1,2}
Time: O(|text|+|pattern|) time; O(|pattern|) auxiliary memory.

```
vector<int> prefixFunction(const string &text) {
    int n = sz(text);
    vector<int> pi(n);
    for (int i = 1; i < n; i++) {
        int j = pi[i - 1];
        while (j > 0 && text[i] != text[j]) j = pi[j - 1];
        if (text[i] == text[j]) j++;
        pi[i] = j;
    }
    return pi;
}
vector<int> kmpSearch(const string &text, const string &pattern
) {
```

```
vector<int> positions;
int patternLength = sz(pattern);
if (patternLength == 0) {
    for (int i = 0; i <= sz(text); i++) positions.push_back(i);
    return positions;
}
vector<int> prefix = prefixFunction(pattern);
int matched = 0;
for (int i = 0; i < sz(text); i++) {
    while (matched > 0 && text[i] != pattern[matched]) matched
    = prefix[matched - 1];
    if (text[i] == pattern[matched]) matched++;
    if (matched == patternLength) {
        positions.push_back(i - patternLength + 1);
        // Keep the longest border so overlapping matches are
        found.
        matched = prefix[matched - 1];
    }
}
return positions;
}
```

Z-Function

Description: z[i] is the common-prefix length of the string and its suffix starting at i; z[0]=0. For matching use pattern+separator+text, with a separator absent from both, and look for z[i]>=pattern.size().
Usage: auto z=zFunction("abacaba"); // {0,0,1,0,3,0,1}
Time: O(N) time and memory.

```
vector<int> zFunction(const string &text) {
    int n = sz(text);
    vector<int> z(n, 0);
    int left = 0;
    int right = -1;
    for (int i = 1; i < n; i++) {
        // [l,r] is the rightmost segment known to match a prefix
        of s.
        if (i <= right) z[i] = min(right - i + 1, z[i - left]);
        while (i + z[i] < n && text[z[i]] == text[i + z[i]]) z[i]
        ]++;
        if (i + z[i] - 1 > right) {
            left = i;
            right = i + z[i] - 1;
        }
    }
    return z;
}
```

Manacher Palindromes

Description: For each position in a string, computes p[0][i] = half length of longest even palindrome around pos i, p[1][i] = longest odd (half rounded down). Indices are 0-based. Odd center i has length 2*p[1][i]+1 and covers [i-p[1][i],i+p[1][i]]. Even center i is the gap before s[i]: length 2*p[0][i], covering [i-p[0][i],i+p[0][i]-1]. The extra even entry at i=n is zero. Radii permit constant-time palindrome checks after construction.

Usage: auto p=manacher("abba");
int evenLength=2*p[0][2]; // 4
Time: O(N)

```
array<vector<int>, 2> manacher(const string &s) {
    int n = sz(s);
    array<vector<int>, 2> radius = {vector<int>(n + 1), vector<
    int>(n)};
    for (int parity = 0; parity < (2); ++parity)
        for (int i = 0, windowLeft = 0, windowRight = 0; i < n; i
        ++){
            int remaining = windowRight - i + !parity;
            if (i < windowRight)
                radius[parity][i] = min(remaining, radius[parity][
                windowLeft + remaining]);
```

```
int left = i - radius[parity][i], right = i + radius[parity][i] - !parity;
while (left >= 1 && right + 1 < n && s[left - 1] == s[right + 1])
    radius[parity][i]++, left--, right++;
if (right > windowRight) windowLeft = left, windowRight = right;
}
return radius;
}
```

Minimum String Rotation

Description: Returns the 0-based start of a lexicographically smallest cyclic rotation, not the rotated string. Periodic input may have equivalent starts; empty input returns 0.
Usage: string s="baca"; int start=minRotation(s); rotate(s.begin(),s.begin()+start,s.end()); // "abac"
Time: $\mathcal{O}(N)$

```
int minRotation(string s) {
    int bestStart = 0, n = sz(s);
    s += s;
    for (int candidate = 0; candidate < (n); ++candidate)
        for (int offset = 0; offset < (n); ++offset) {
            if (bestStart + offset == candidate || s[bestStart + offset] < s[candidate + offset]) {
                candidate += max(0, offset - 1);
                break;
            }
            if (s[bestStart + offset] > s[candidate + offset]) {
                bestStart = candidate;
                break;
            }
        }
    return bestStart;
}
```

Rolling Hash

Description: Polynomial prefix hash with base 31 and modulus 1000000003. getHash uses 1-based inclusive [l,r], although the input string is normal 0-based. Compare equal-length substrings. A single fixed hash can collide: use KMP for exact matching, or two independent moduli when adversarial collisions matter. No updates; this version expects lowercase a.z.
Usage: StringHash hash("ababa");
bool same=hash.getHash(1,3)==hash.getHash(3,5); // both "aba"
StringHash pattern("aba");
bool match=hash.getHash(1,3)==pattern.getHash(1,3);
Time: $\mathcal{O}(N)$ construction and memory; $\mathcal{O}(1)$ per substring hash.

```
struct StringHash {
    static constexpr ll base = 31, mod = 1000000003;
    vector<ll> power, prefix;
    StringHash(const string &text) : power(sz(text) + 1), prefix(sz(text) + 1) {
        power[0] = 1;
        for (int i = 1; i <= sz(text); i++) {
            power[i] = power[i - 1] * base % mod;
            prefix[i] = (prefix[i - 1] * base + text[i - 1] - 'a' + 1) % mod;
        }
    }
    ll getHash(int l, int r) const {
        return (prefix[r] - prefix[l - 1] * power[r - l + 1] % mod + mod) % mod;
    }
};
```

Various (10)

10.1 Intervals

Interval Union and Removal

Description: Add and remove intervals from a set of disjoint intervals. Will merge the added interval with any overlapping intervals in the set when adding. Intervals are [inclusive, exclusive). Maintain covered integer-coordinate ranges rather than individual points. Touching intervals merge too: [1,3) and [3,5) become [1,5). removeInterval can split an interval. Use $l \leq r$; empty intervals have no effect. A call touching k stored intervals costs $\mathcal{O}(\log N + k)$, because those intervals must be erased.
Usage: set<pii> intervals; addInterval(intervals,1,5); removeInterval(intervals,2,4); // {[1,2),[4,5)}
Time: $\mathcal{O}(\log N)$ amortized; $\mathcal{O}(\log N + k)$ for a call erasing k intervals.

```
set<pii>::iterator addInterval(set<pii> &intervals, int l, int r) {
    if (l == r) return intervals.end();
    auto it = intervals.lower_bound({l, r}), hint = it;
    while (it != intervals.end() && it->first <= r) {
        r = max(r, it->second);
        hint = it = intervals.erase(it);
    }
    if (it != intervals.begin() && (--it)->second >= l) {
        l = min(l, it->first);
        r = max(r, it->second);
        intervals.erase(it);
    }
    return intervals.insert(hint, {l, r});
}

void removeInterval(set<pii> &intervals, int l, int r) {
    if (l == r) return;
    auto it = addInterval(intervals, l, r);
    auto oldRight = it->second;
    if (it->first == l)
        intervals.erase(it);
    else
        (int &)it->second = l;
    if (r != oldRight) intervals.emplace(r, oldRight);
}
```

Minimum Interval Cover

Description: Compute indices of smallest set of intervals covering another interval. Intervals should be [inclusive, exclusive). To support [inclusive, inclusive], change (A) to add || chosen.empty(). Returns empty set on failure (or if G is empty). A greedy sweep chooses the eligible interval reaching farthest right. Output is 0-based indices into the original input, not its sorted order. Minimizes interval count rather than total cost. Adjacent intervals may meet at an endpoint; a gap causes failure. Check whether the target is empty before interpreting an empty result as failure.
Usage: vector<pii> intervals={{0,3},{2,5},{0,1}}; auto chosen=cover(pii(0,5),intervals); // {0,1}
Time: $\mathcal{O}(N \log N)$

```
template <class T> vector<int> cover(pair<T, T> target, vector<pair<T, T>> intervals) {
    vector<int> order(sz(intervals)), chosen;
    iota(all(order), 0);
    sort(all(order), [&](int a, int b) { return intervals[a] < intervals[b]; });
    T coveredTo = target.first;
    int nextInterval = 0;
    while (coveredTo < target.second) { // (A)
        pair<T, int> farthest = make_pair(coveredTo, -1);
        while (nextInterval < sz(intervals) && intervals[order[nextInterval]].first <= coveredTo) {
            farthest =
                max(farthest, make_pair(intervals[order[nextInterval]].second, order[nextInterval]));
        }
    }
```

```
        nextInterval++;
    }
    if (farthest.second == -1) return {};
    coveredTo = farthest.first;
    chosen.push_back(farthest.second);
}
return chosen;
}
```

Constant-Value Interval Search

Description: Split a monotone function on [from, to) into a minimal set of half-open intervals on which it has the same value. Runs a callback g for each such interval. Use when evaluating a monotone integer-indexed function is expensive and it changes value only a few times. visitInterval receives [lo,hi,value) for each maximal [lo,hi). Increasing/decreasing functions both work. A value must not disappear then return, or equal endpoints can hide changes. valueAt must be deterministic.
Usage: constantIntervals(0,10,[](int i) { return i/3; }, [](int lo,int hi,int value) { ... });
Time: $\mathcal{O}(k \log \frac{n}{k})$

```
template <class F, class G, class T>
void rec(int from, int to, F &valueAt, G &visitInterval, int &intervalStart, T &currentValue, T endValue) {
    if (currentValue == endValue) return;
    if (from == to) {
        visitInterval(intervalStart, to, currentValue);
        intervalStart = to;
        currentValue = endValue;
    } else {
        int mid = (from + to) >> 1;
        rec(from, mid, valueAt, visitInterval, intervalStart, currentValue, valueAt(mid));
        rec(mid + 1, to, valueAt, visitInterval, intervalStart, currentValue, endValue);
    }
}

template <class F, class G> void constantIntervals(int from, int to, F valueAt, G visitInterval) {
    if (to <= from) return;
    int intervalStart = from;
    auto currentValue = valueAt(intervalStart), endValue = valueAt(to - 1);
    rec(from, to - 1, valueAt, visitInterval, intervalStart, currentValue, endValue);
    visitInterval(intervalStart, to, endValue);
}
```

10.2 Misc. algorithms

Ternary Search

Description: Returns the smallest maximizing integer index on inclusive [lo,hi]. Require a unimodal function that rises strictly before its maximum plateau, then never rises again. To minimize, reverse both comparisons marked (A) and (B).
Usage: int best=ternSearch(0,10,[](int i) { return -(i-4)*(i-4); });
// best=4
Time: $\mathcal{O}(\log(b - a))$

```
template <class F> int ternSearch(int lo, int hi, F valueAt) {
    assert(lo <= hi);
    while (hi - lo >= 5) {
        int mid = (lo + hi) / 2;
        if (valueAt(mid) < valueAt(mid + 1))
            lo = mid; // (A)
        else
            hi = mid + 1;
    }
}
```

```
for (int i = lo + 1; i < (hi + 1); ++i)
    if (valueAt(lo) < valueAt(i)) lo = i; // (B)
return lo;
}
```

Fast Subset Sum (Knapsack)

Description: Given N non-negative integer weights w and a non-negative target t, computes the maximum $S \leq t$ such that S is the sum of some subset of the weights. This is 0/1 subset sum with bounded weights: each weight is used at most once and value equals weight. Returns only the best achievable sum, not chosen items. Weights/capacity must be nonnegative and sums must fit int. Useful when maxWeight is small even if capacity is large. This is not general profit/weight knapsack.

Usage: int best=knapsack({3,5,7},10); // 10, using 3+7
Time: $\mathcal{O}(N \max(w_i))$

```
int knapsack(vector<int> weights, int capacity) {
    int total = 0, prefixSize = 0, state;
    while (prefixSize < sz(weights) && total + weights[prefixSize]
        ] <= capacity)
        total += weights[prefixSize++];
    if (prefixSize == sz(weights)) return total;
    int maxWeight = *max_element(all(weights));
    vector<int> previous, dp(2 * maxWeight, -1);
    dp[total + maxWeight - capacity] = prefixSize;
    for (int i = prefixSize; i < (sz(weights)); ++i) {
        previous = dp;
        for (int state = 0; state < (maxWeight); ++state)
            dp[state + weights[i]] = max(dp[state + weights[i]],
                previous[state]);
        for (state = 2 * maxWeight; --state > maxWeight;)
            for (int j = max(0, previous[state]); j < (dp[state]); ++j)
                dp[state - weights[j]] = max(dp[state - weights[j]], j)
                ;
    }
    for (total = capacity; dp[total + maxWeight - capacity] < 0;
        total--);
    return total;
}
```

10.3 Dynamic programming

SOS DP: Subsets and Supersets

Description: In-place sums over subsets/supersets, including the mask it-self and zero. dp has exactly 2^B entries ($B \geq 0$); copy the input separately for each transform. For frequencies, sub[full XOR mask] counts disjoint input masks. inverse subtracts to restore input. Keep the bit loop outermost. Intermediate sums must fit ll; normalize modular results. Min/max variants have no subtraction inverse.

Usage: vector<ll> freq={1,2,3,4}; // masks 0,1,2,3
auto sub=freq; subsetSOS(sub); // {1,3,4,10}
auto super=freq; supersetSOS(super); // {10,6,7,4}
int full=sz(freq)-1, mask=1;
ll disjoint=sub[full^mask]; // 4
subsetSOS(sub,true); // restores freq

Time: $\mathcal{O}(B2^B)$ time; $\mathcal{O}(1)$ auxiliary memory.

```
void subsetSOS(vector<ll> &dp, bool inverse = false) {
    int bits = __lg(sz(dp));
    for (int bit = 0; bit < bits; bit++)
        for (int mask = 0; mask < sz(dp); mask++)
            if (mask >> bit & 1) dp[mask] += (inverse ? -1 : 1) * dp[
                mask ^ (1 << bit)];
}
void supersetSOS(vector<ll> &dp, bool inverse = false) {
    int bits = __lg(sz(dp));
    for (int bit = 0; bit < bits; bit++)
        for (int mask = 0; mask < sz(dp); mask++)
```

```
if (!(mask >> bit & 1)) dp[mask] += (inverse ? -1 : 1) *
    dp[mask | (1 << bit)];
}
```

Knuth DP Optimization

Description: When doing DP on intervals: $a[i][j] = \min_{i < k < j} (a[i][k] + a[k][j]) + f(i, j)$, where the (minimal) optimal k increases with both i and j , one can solve intervals in increasing order of length, and search $k = p[i][j]$ for $a[i][j]$ only between $p[i][j - 1]$ and $p[i + 1][j]$. This is known as Knuth DP. Sufficient criteria for this are if $f(b, c) \leq f(a, d)$ and $f(a, c) + f(b, d) \leq f(a, d) + f(b, c)$ for all $a \leq b \leq c \leq d$. Consider also: CHT (ch. Data structures), monotone queues, ternary search. This is a re-currence recipe, not a callable function. Typical use is optimal merging of adjacent segments with f(i,j) equal to interval sum. Initialize base intervals and optimal split positions, then process increasing lengths. The required monotonicity concerns optimal splits, not DP values; prove it or establish the sufficient inequalities before using the bound.

Usage: // For each increasing interval [i,j], try only splits
// opt[i][j-1] <= k <= opt[i+1][j], also restricted to i < k < j.
// Save dp[i][j] and its smallest minimizing split.
Time: $\mathcal{O}(N^2)$

Divide-and-Conquer DP

Description: Given $a[i] = \min_{lo(i) \leq k < hi(i)} (f(i, k))$ where the (minimal) optimal k increases with i , computes $a[i]$ for $i = L..R - 1$. Customiza-tion skeleton for one DP layer. Replace lo/hi with valid split bounds, f with previousLayer[split]+cost(split,index), and store with destination ar-rays. solve(l,r) covers [l,r). Each index needs at least one candidate; optimal split indices must be nondecreasing. Keep the previous layer fixed while computing the current one.

Usage: // Implement lo,hi,f,store and declare DP arrays first.
DivideConquerDP solver; solver.solve(l,n+1); // indices 1..n
Time: $\mathcal{O}((N + (hi - lo)) \log N)$

```
struct DivideConquerDP { // Modify at will:
    int lo(int index) { return 0; }
    int hi(int index) { return index; }
    ll f(int index, int split) { return dp[index][split]; }
    void store(int index, int split, ll value) { answer[index] =
        pii(split, value); }
    void rec(int l, int r, int optL, int optR) {
        if (l >= r) return;
        int mid = (l + r) >> 1;
        pair<ll, int> best(LLONG_MAX, optL);
        for (int split = max(optL, lo(mid)); split < (min(optR, hi(
            mid))); ++split)
            best = min(best, make_pair(f(mid, split), split));
        store(mid, best.second, best.first);
        rec(l, mid, optL, best.second + 1);
        rec(mid + 1, r, best.second, optR);
    }
    void solve(int l, int r) { rec(l, r, INT_MIN, INT_MAX); }
};
```

10.4 Optimization tricks

__builtin_ia32_ldmxcsr(40896); disables denormals
(which make floats 20x slower near their minimum value).

10.4.1 Bit hacks

- $x \ \& \ -x$ is the least bit in x.
- for (int x = m; x;) { --x &= m; ... } loops over all subset masks of m (except m itself).
- $c = x \ \& \ -x, \ r = x + c; ((r \wedge x) \gg 2) / c \mid r$ is the next number after x with the same number of bits set.

10.4.2 Pragmas

- **#pragma** GCC optimize ("Ofast") will make GCC auto-vectorize loops and optimizes floating points better.
- **#pragma** GCC target ("avx2") can double performance of vectorized code, but causes crashes on old machines.
- **#pragma** GCC optimize ("trapv") kills the program on integer overflows (but is really slow).